

AI is done chatting, now give it tasks

For years, AI was something you opened when you needed an answer. You asked it to explain a term, rewrite a paragraph, summarise a document, generate a caption or settle a friendly argument. That was impressive enough. But it was still mostly a smarter search box with better manners. The next phase is different. AI is moving from answering questions to doing work.

That shift is already visible in the numbers. ChatGPT crossing one billion monthly active app users is not just another Silicon Valley milestone; it is a sign that AI has become a mass habit. Globally, 88 percent of surveyed organisations now report regular AI use in at least one business function. Even more interestingly, 62 percent are already experimenting with AI agents, while 23 percent say they are scaling agentic systems somewhere inside the enterprise. In India, the timing could not be more important: the country has 958 million active internet users, 57 percent of them in rural India, and 44 percent of users have already engaged with AI-enabled features such as voice search, image search, chatbots or AI filters. Meanwhile, the IndiaAI Mission has onboarded more than 38,000 GPUs for common compute access.

In other words, AI agents are not waiting politely at the edge of the future. They are entering the workplace, classroom, phone, browser, cloud drive, CRM dashboard, coding editor, customer-support queue and family planning spreadsheet right now.

But here is the catch: most people are still using AI like it is 2023. They type a vague prompt, admire or dislike the answer, and move on. Agents demand a different habit. You do not merely ask, "Can you help me?" You define a task, share context, set limits, choose tools, demand verification and decide where the system must stop for approval. That is the difference between playing with AI and working with AI.

This *Fast Track* is built around that difference. It is not about treating AI agents as magical employees that will run your life while you drink coffee. That fantasy is both lazy and dangerous. A badly briefed agent can misunderstand, hallucinate, expose private data, send the wrong message, overwrite the wrong file or confidently create a mess at machine speed. McKinsey's 2025

survey found that 51 percent of organisations using AI had already seen at least one negative consequence, with inaccuracy among the major concerns. The lesson is clear: power without guardrails is not productivity, but risk in a nicer interface.

The smarter approach is supervised delegation. Let agents summarise the boring report, but check the conclusion. Let them draft the client email, but approve it yourself. Let them compare laptops, policies or vendors, but verify current prices and fine print. Let them organise notes, create revision plans, classify leads, debug code, prepare content calendars and build

simple workflows. Keep judgement, ethics, privacy and final decisions with the human.

For Indian readers, this is especially practical. A student can build a revision agent without outsourcing learning. A teacher can create worksheets in English and Hindi without losing classroom control. A freelancer can track proposals, invoices and follow-ups without hiring an assistant. A small business can sort leads and draft replies without sounding like a bot. A journalist can turn transcripts into quote banks and story angles. A developer can use coding agents without blindly shipping insecure code.

The next wave of digital literacy will not be about knowing every tool name. Those will change. It will be about knowing how to brief, connect, restrict, review and improve agents. If you can do that, AI stops being a novelty and becomes infrastructure. **d**

