

A useful rule is to make an adjustment, close your eyes and toggle it on and off. If “processed” merely sounds different rather than clearly better, use less.

Compression: Make voices easier to follow

Podcast speech can move from a quiet sentence to loud laughter within seconds. Compression reduces that dynamic range by turning down signal above a chosen threshold. Adobe describes compressor controls such as threshold, ratio, attack and release, with make-up gain available to restore overall level afterwards.

For beginners, start gently. A ratio around 2:1 or 3:1 can be a reasonable learning point for dialogue, but treat it as an experiment rather than a rule.

- **Step 1:** Set compression while listening to an energetic part of the recording.
- **Step 2:** Lower the threshold until louder phrases trigger moderate gain reduction.
- **Step 3:** Compare processed and unprocessed audio at similar perceived volume.
- **Step 4:** Stop if breaths, room noise and mouth sounds begin jumping forward unnaturally.

Compression should make the presenter easier to follow, not make every syllable equally loud.

Fix harsh “S” sounds

Some voices and microphones produce exaggerated sibilance: sharp “s”, “sh” and “ch” sounds that become tiring on headphones.

A de-esser reduces energy in the frequency region where that harshness occurs. Adobe Audition and Fairlight both provide dedicated de-essing tools.

Use the smallest amount that solves the problem. Too much de-essing can make the speaker sound as if they have developed a lisp. If only three words are offensive, manually reducing those moments may sound more natural than aggressively processing the entire recording.

Balance multiple speakers



PERFORM THE THREE-DEVICE TEST



Before publishing, listen to representative sections on three devices:

- **Step 1:** Good headphones — to hear clicks, hiss and editing artifacts.
- **Step 2:** Ordinary earbuds — to approximate common listening conditions.
- **Step 3:** A phone speaker — to expose speech that depends too heavily on bass or has music masking the voice.

Check the opening, the loudest section, a quiet section, every inserted advertisement or clip, and the final minute.

Then check the file name, episode number and metadata before uploading.

Once individual voices are cleaned up, make them feel as though they belong in the same programme.

Work through the conversation and adjust clip or track gain so listeners do not need to reach for the volume control whenever the speaker changes. Use short gain adjustments when somebody suddenly moves away from the microphone rather than compressing the entire track excessively.

If two microphones captured the same person, mute unwanted bleed wherever practical, but avoid edits that make the background ambience jump unnaturally.

Add music without burying speech

Music should perform a job: establish identity, signal a transition, create atmosphere or provide an ending. It should not run underneath the entire episode simply because you found a track you like.

Place intro music on a separate track. Fade it down as dialogue begins. Bring it back only where it contributes something.

This technique is often called ducking. DaVinci Resolve’s Fairlight tools, for example, can automatically lower a music track when dialogue is present and restore it afterwards.