

ChatGPT 5

The DOs and DON'Ts of AI training according to OpenAI

—By **Satvik Pandey**

Artificial Intelligence is no longer the distant vision of futurists – it is here, embedded in our daily lives, shaping how we work, interact, and even make decisions. But with this growing integration comes an undeniable responsibility – ensuring that AI systems behave in ways that are helpful, safe, and aligned with human values. OpenAI's Model Spec represents a major step toward defining and regulating this behavior. While technical in its scope, its implications extend far beyond developers and policymakers; it touches every user who engages with AI.

Why AI needs a behavioral blueprint?

For many, AI governance might seem like an abstract concept, best left to engineers and ethicists. But consider this: each time you use an AI-powered assistant, engage with an AI-curated feed, or rely on AI for creative or professional work, you are interacting with a system designed to follow certain behavioral principles. These principles dictate how an AI should respond to your queries, handle ambiguous requests, and navigate ethical dilemmas.

OpenAI's Model Spec outlines a structured approach to aligning AI behavior with human expectations. The document is built on three fundamental objectives: maximising user benefit, minimising harm, and ensuring that AI operations remain sustainable within legal and ethical boundaries. This trinity of principles functions as a guiding framework, balancing the trade-offs inherent in AI decision-making.

Structuring AI decision-making

One of the most fascinating aspects of the Model Spec is how it structures

decision-making through a layered chain of authority. The AI system adheres to a hierarchy that starts with platform-level rules – strict, non-negotiable directives aimed at preventing harm and maintaining compliance with laws. Beneath that are developer-imposed guidelines, offering flexibility for customisation, and finally, user-level instructions, allowing individuals to tailor AI behavior within predefined boundaries.



For example, an AI model might be designed to summarise news articles while avoiding the spread of misinformation. At a platform level, it would be restricted from generating false claims, even if a developer were to tweak it for a specific use case. At the developer level, adjustments could be made to fine-tune how summaries are generated. At the user level, personal preferences – such as a preference for concise versus detailed summaries – could be applied.

This hierarchical approach ensures that AI does not operate in a vacuum. Instead, it dynamically adapts while maintaining a core set of ethical constraints.

How the Model Spec will shape ChatGPT-5?

The release of ChatGPT-5 is expected to mark a significant leap in AI capabilities, and the Model Spec will serve as the foundation for its development. OpenAI

will use this framework to fine-tune the model, ensuring that it adheres to ethical guidelines while providing more accurate and context-aware responses.

One key area of improvement is adaptive alignment, where ChatGPT-5 will better interpret nuanced user intent while staying within ethical constraints. This means a model that is more responsive, yet still bound by rules that prevent harmful or misleading outputs.

Additionally, OpenAI plans to use the Model Spec to refine self-supervised learning, where the AI continuously evaluates its own responses against predefined safety and accuracy measures. This could reduce the instances of misinformation and help mitigate bias in AI-generated content. However, despite these precautions, aligning AI with human values is not a perfect science. And that leads to a pressing concern – what happens when AI models trained under these principles become too restrictive or, conversely, too lenient?

When AI alignment leads to rogue models?

While the Model Spec aims to create safer and more aligned AI, critics argue that such frameworks could inadvertently lead to the development of rogue AI models – either due to over-regulation or deliberate misuse.

The risk of over-alignment

AI models trained with excessive constraints may become frustratingly restrictive, refusing to engage with complex or controversial topics, even when it would be reasonable to do so. Over time, this could drive users toward alternative AI systems with fewer restrictions – potentially leading to the rise of unregulated, uncensored models trained outside of ethical AI frameworks. **[d]**

