


OpenAI first launched ChatGPT Plus at \$20 a month, then an enterprise version at \$200 per month! In this scenario, the end-user is always paying more and more for incremental improvements. And it's not just OpenAI. A handful of well-funded AI startups have banked on the narrative that advanced AI is too complicated for any average Joe to replicate cheaply. That sentiment, ironically, is losing credence fast.

DeepSeek R1's training budget was only \$5 million compared to \$7 billion by OpenAI for ChatGPT training and inference costs in 2024. If that's even half true, it signals a tectonic shift. Maybe, just maybe advanced AI doesn't need the mountain of GPUs we thought. Or at least not as many.

That's not me cheerleading for someone's downfall, it's just me observing that maybe we never fully knew how resource-light advanced model training can become.

But then DeepSeek R1 waltzes in with a price point that's 20 to 30 times cheaper than the major US players. Of course, that's no small change, enough for big enterprise customers to start wondering if they can get 90% of the top-tier AI performance from an open-source or far cheaper model? It's reminiscent of last summer's open-source wave, when Meta released Llama, or how stable diffusion-based AI image models became the open alternative to certain big brand commercial solutions. DeepSeek's approach is a new iteration of that same spirit – bring robust AI to the masses at minimal cost, shake up the status quo.

Another dimension to consider here is the potential slowdown in



AI would replace every writer or coder in the next 12 months? We survived that wave of doomsday predictions. Now we're hearing a parallel about how big AI labs might topple if open-source or cost-efficient solutions overshadow them.

So, does DeepSeek's arrival mean we're about to see OpenAI vanish into irrelevance? Probably not. OpenAI's name recognition and the quality of ChatGPT remain potent. But it does highlight an undercurrent that's impossible to ignore and will continue to grow – that open-source or budget AI solutions can be shockingly good, if done right. Large enterprise customers might continue paying for top-tier GPT-like reliability, while smaller players lean on open solutions like DeepSeek. The AI marketplace may just get

more diverse.

For me, the sweet spot is an ecosystem where major players push the envelope in specialised domains, while open-source alternatives keep them honest on pricing and accessibility. Historically, that's how software evolves. Look at Linux vs. Windows or Android vs. iOS. We see an interplay of commercial scale and open, community-based innovation. In a best-case scenario, we get more consumer choice and cheaper or even free AI offerings.

At the end of the day, free or cheap AI doesn't necessarily undermine the viability of well-funded AI startups. It simply forces them to become better and more transparent about what exactly we're paying for. So, ironically, maybe we owe DeepSeek a thank-you, not just for unveiling some cheaper approach but for reminding us all that innovation doesn't have to come with an overblown price tag. **d**

NVIDIA revenue. NVIDIA is the biggest winner of the AI hardware gold rush. If AI developers pivot to smaller, more efficient training regimes – like how DeepSeek supposedly managed with \$5 million in compute – NVIDIA might not see that monstrous GPU demand continue at its current breakneck pace.

If these open-source, cost-efficient solutions become widely adopted, what does it mean for giant AI startups flush with billions of dollars in VC money? Or for NVIDIA, who's made a fortune selling advanced GPUs to train these gargantuan neural nets? Potentially, a large chunk of that profit model gets disrupted if the rest of the industry follows DeepSeek's lightweight, minimal hardware approach.

What does 'DeepSeek moment of AI' mean for us

Remember the headlines from a year ago, proclaiming how Generative