




NVIDIA DGX Spark lights up AI supercomputing

Rebranded DIGITS debuts with Blackwell DGX Station for desktop AI power

Mithun Mohandas | mithun@digit.in

NVIDIA has introduced two new personal AI supercomputers, the DGX Spark and DGX Station, as part of its efforts to make high-performance AI computing more accessible. These systems, powered by the Grace Blackwell platform, are designed to cater to AI developers, researchers, data scientists, and students who require powerful computing capabilities for prototyping, fine-tuning, and running AI models. The NVIDIA DGX Spark will be priced at USD 3000, but there's no price for the DGX Station yet. We have seen previous iterations of the DGX Station start from USD 99,000. However, while NVIDIA presents these as entirely new AI

computing solutions, DGX Spark is, in reality, a rebranded and updated version of Project DIGITS. DIGITS was originally launched in 2015 as an AI development platform aimed at deep learning practitioners. It was built on top of CUDA and cuDNN, making it essentially a middleware for deep learning GPU training.



DIGITS would be packaged with DRIVE as part of NVIDIA's training solution for designing self-driving cars. DIGITS was first introduced back in March 2015 at NVIDIA's GPU Technology Conference. Since then NVIDIA has been iterating on the software with version 2 of the package. Back then DIGITS was NVIDIA's higher-level neural network software for general scientists and researchers (as opposed to programmers). Earlier this year at CES, NVIDIA announced that it would be packaging DIGITS with a GPU making it an out-of-the-box solution for anyone to build a neural network training system.

Call it the Mac Mini of AI training, if you will. With the introduction of the Blackwell GPU architecture, NVIDIA has now repositioned this product under the DGX branding, emphasizing its suitability for AI-native applications.

DGX Spark: The Mac Mini of AI Supercomputers

NVIDIA positions DGX Spark as the smallest AI supercomputer, bringing powerful AI capabilities to desktop environments. At its core is the NVIDIA GB10 Grace Blackwell Superchip, featuring a Blackwell GPU with fifth-generation Tensor Cores and FP4 precision support.

This combination enables the system to deliver up to 1,000 TOPS (trillion operations per second) of AI compute, making it well-suited for fine-tuning and inference tasks involving modern AI models such as NVIDIA's Cosmos Reason and GR00T N1 foundation models.

A key feature of the GB10 Superchip is its use of NVIDIA NVLink-C2C interconnect technology. This provides a CPU-GPU coherent memory model with five times the bandwidth of PCIe 5.0, significantly improving memory access speeds for AI workloads. This architecture ensures that data-intensive applications, including large-scale generative AI