

Intel's India vision

From AI acceleration to enabling electronics manufacturing for the world

Jayesh Shinde | jayesh@digit.in

It's hard not to notice India's steady rise in the global technology conversation – be it in AI, semiconductor manufacturing, or data center innovation, according to Santhosh Viswanathan, Vice President & Managing Director, Intel, India Region, who has a ringside seat to all the cutting-edge action and raw dynamism shaping India's tech future.

In an extensive interview, he discussed Intel's AI hardware approach, the "Make in India" initiative, and insights into some distinct advantages of India both as a massive consumer market and a hub for future-looking R&D. Santhosh Viswanathan's perspective sees India not merely as a place to do business but as a crucible for invention – especially in AI, where more than one-fifth of the world's data is generated. And according to him, integrating all that data with real-world solutions, though, demands a flexible approach. For Intel, that means avoiding the trap of "AI = GPU," a phrase Santhosh is quick to dispel.



Santhosh Viswanathan
Vice President & Managing Director,
Intel, India Region

Shattering the "AI = GPU" myth

"Intel recognizes that AI solutions must be tailored to meet the unique needs of different markets," Santhosh says, highlighting how crucial it is to sidestep the assumption that AI is always about a GPU. The company's strategy, he explains, hinges on one overwhelming truth, which is that AI can run on many kinds of hardware.

"For instance, Xeon processors are designed to handle a wide range of AI workloads, from edge to cloud, while Gaudi AI chips offer cost-effective AI acceleration for training and inference tasks."

But what about GPUs and large language models (LLMs)? NVIDIA

has become the darling of AI hype in recent years, leaving many to wonder how Intel competes or even coexists in that space. Santhosh acknowledges the heavy odds at stake but offers an alternative that's not just necessary but demanded by tech companies. "AI is not a one-size-fits-all solution. For many, AI is just another workload running alongside hundreds of others. Sometimes you don't need the highest performance on a single function – just enough to meet the SLA and run efficiently." He points out that many businesses use LLMs not to train from scratch, but to refine or fine-tune existing models – tasks well within the capability of a CPU, especially a modern Xeon processor.

So, where do Gaudi chips come in? "Our Gaudi AI chips provide competitive performance and cost-efficiency for large-scale AI training and inference, making them an attractive option for enterprises looking to scale AI deployments," he says. At the same time, he doesn't deny that smaller or medium-sized models can be effectively managed on general-purpose Xeons. The overarching theme is that Intel wants to offer a continuum of solutions, not just a single one.

That's the mantra he keeps returning to – "AI does not equal GPUs." In many enterprise scenarios, a CPU might be more than enough to handle fine-tuning a smaller large language model or to run certain machine learning tasks. Relying solely on a GPU for every job can be overkill – and an expensive one at that. The future, from Intel's vantage point, is about heterogeneous computing: a CPU here, an accelerator there, and a flexible approach that ensures you're using just enough horsepower to get the job done efficiently.

"Our recently launched Intel Xeon 6 processors with Performance-cores, for instance, allow for 5:1 consolidation of a five-year-old server on average, with potential for up to 10:1 in certain use cases," Santhosh