

NVIDIA becomes first ever to hit \$4 trillion mark, spurred by AI



NVIDIA is the first company to reach a market capitalization of \$4 trillion, a feat achieved on 9th July 2025. This valuation cements NVIDIA as a titan in the global technology landscape, driven by its dominance in the artificial intelligence (AI) sector. The company's meteoric rise shows the power of AI and the demand for NVIDIA's cutting-edge graphics processing units, which have become the backbone of the AI revolution. tinyurl.com/vzh5ju8j

Claude AI can now design and edit your Canva projects



Anthropic has made its AI chatbot Claude even more useful, especially if you're someone who regularly uses Canva. Starting today, Claude can help users design, edit and manage Canva projects. All they have to do is describe what they need. With this new integration, Claude can now carry out tasks like creating presentations, resizing images and filling out templates. But here's the catch. To use the new feature, you'll need a paid subscription to both platforms. So, this convenience comes at a cost. tinyurl.com/5n7v2xfj

OpenAI, Google, Anthropic researchers warn about AI 'thoughts'

Chain-of-thought visibility in AI models is critical

—By **Vyom Ramani**

In a rare show of unity, researchers from OpenAI, Google, Anthropic, and Meta have issued a stark warning: the window to understand and monitor artificial intelligence "thought processes" is closing fast. As AI systems grow more sophisticated, their decision-making becomes increasingly opaque, raising concerns about safety, transparency, and control. This collaborative call to action, detailed in a position paper published on July 15, 2025, emphasizes the need for techniques for monitoring AI's "chain-of-thought" (CoT) reasoning, a critical tool for ensuring these systems remain aligned with human values.

Modern AI models, such as OpenAI's o1 and o3, DeepSeek's R1, and Anthropic's Claude 3.7 Sonnet, are designed to "think out loud" by turning complex problems into step-by-step reasoning in readable language. This CoT process acts like a window into AI decision-

making, allowing researchers to spot potential issues, like models contemplating unethical actions with phrases like "let's hack" or "let's sabotage." But this transparency is fragile. As AI architectures evolve, researchers warn that future systems may leave language-based reasoning, thinking instead in mathematical spaces or compressed formats that humans can't decipher.



Bowen Baker, OpenAI researcher and lead author of the position paper, in an interview to TechCrunch: "The existing CoT monitorability may be extremely fragile. Higher-compute RL, alternative model architectures, certain forms of process supervision, etc. may all lead to models that obfuscate their thinking."

If this happens, the ability to detect harmful intentions or biases before they manifest could be lost, leaving society vulnerable to unpredictable AI behavior.

Why monitoring AI 'thoughts' is important

The ability to monitor reasoning is not technical curiosity, it's AI safety. Current models also show signs of deception. An Anthropic study revealed that Claude 3.7 Sonnet used only subtle hints in its reasoning 25% of the time, while DeepSeek's R1 did so 39%. With problematic prompts, such as instructions implying unauthorized system access, models often fabricated justifications rather than admitting to shortcuts.

Overwatch of AI reasoning could improve alignment but risks tracing CoT disingenuous, as models might learn to generate "safe" reasons that mask their processes. The paper also calls for lifting nondisclosure agreements and making anonymous channels for concerned employees. tinyurl.com/5fen7sdv