Apple plans AI comeback with tabletop robot, lifelike Siri

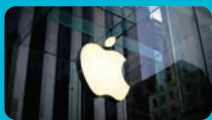

Apple is reportedly preparing a major push into artificial intelligence with a range of new devices, including a lifelike version of Siri, a smart speaker with a display, and home security cameras. According to Bloomberg, these products are Apple's effort to comeback in the AI race. Their centrepiece is a tabletop robot, likely featuring an iPad-like display on a movable limb and run on a new version of Siri that remembers details and interacts naturally. **tinyurl.com/fh7mf8xn**

Is Google's AI Search killing website traffic?

There's been a lot of talk lately about whether Google's new AI-powered Search experience is taking away traffic from websites. With features like AI Overviews and AI Mode, people can now ask longer, more detailed questions and get quick summaries directly in Search. This has raised concerns over declining website traffic which Google addressed in a recent blogpost. They claim that website traffic from Google Search has remained "relatively stable" year-over-year, despite what some third-party reports suggest. **tinyurl.com/y4safu57**

Persona Vectors: Anthropic's solution to monitor and control AI behaviour

New AI method steers model behavior without retraining or prompt engineering

—By **Vyom Ramani**

I've chatted with enough bots to know when something feels a little off. Sometimes, they're overly flattering. Other times, weirdly evasive. And occasionally, they take a hard left into completely bizarre territory. So when Anthropic dropped its latest research on "Persona Vectors" – a technique to understand and steer a model's behavior without re-training it – I knew this was more than another AI safety buzzword.

How do Persona Vectors work?

Persona vectors are internal activation patterns inside AI models that correspond to specific "traits" like sycophancy, hallucination, or maliciousness. Anthropic's researchers found that when a model consistently behaves a certain way, say, by excessively flattering the user, that behavior creates a measurable pattern in the model's neural activations.

By comparing these patterns to those from neutral behavior, they isolate a vector – essentially a direction in the model's internal space – that represents that trait. During inference, developers can inject this vector to amplify the behavior or subtract it to suppress it.

In practice, this opens up new ways to control model behavior. If a chatbot is too much of a people-pleaser, subtracting the sycophancy vector can make it more assertive. If it tends to hallucinate, steering away from the hallucination vector makes it more cautious.

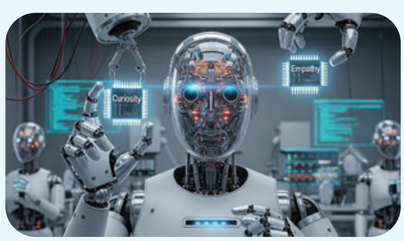

Anthropic also uses persona vectors during fine-tuning in a process they call preventative steering. Here, they deliberately inject harmful traits like "evil" into the model during training, not to corrupt it, but to build resistance. Inspired by the concept of vaccines, this helps the model learn to ignore or reject bad behavior patterns later on. Importantly, these harmful vectors are disabled at deployment, so the final model behaves as intended but is stable and aligned.

Does it actually work?

Yes, it works and has been tested across multiple open-source models like Qwen 2.5 and Llama 3.1. Injecting or removing vectors consistently altered behavior without damaging core performance. And when applied during fine-tuning, the models became more resistant to adopting harmful traits later.

Even better, benchmark scores (like MMLU) stayed strong. That means you don't lose intelligence by improving alignment. Traditionally, controlling AI behavior meant either prompt engineering or retraining the whole model. Persona vectors offer a third path: precise, explainable, and fast.
tinyurl.com/mr3uxtsw