

Open models key for AI's future

NVIDIA's Kari Briski thinks it will define the next era of AI

Jayesh Shinde | editor@digit.in

In the world of computing, open (be it software or hardware) is a curious and oft-abused word. Thrown around left, right and centre, but rarely examined. Some say open means freedom from corporate control. Others use it as shorthand for transparency or accessibility. But for Kari Briski, Vice President of Generative AI Software for Enterprise at NVIDIA, who leads NVIDIA's Nemotron model family and the ecosystem of software around it, "open" isn't just a slogan.

"We believe that to further AI development you have to have open models," Briskii tells me, while referring to NVIDIA's Nemotron model family. Nemotron constitutes AI models and techniques that are developed by NVIDIA to build advanced reasoning-based AI agents. "We do it for ourselves, and are talking about it now so people understand this is like CUDA – like our math libraries – you can depend on it. We are committed to it. It is part of our strategy because first and foremost we build it for ourselves, and secondly because we believe in this AI development."

In other words, Nemotron is part of NVIDIA's core strategy. And it

arrives right as enterprises move from sandbox GenAI experiments to shipping systems at scale.

Briski is clear-eyed about the industry's semantic slop around open source software,



Kari Briski
VP – GenAI Software for Enterprise, NVIDIA

especially in the uber competitive GenAI landscape. "When you talk about open models, there's been an explosion since Llama 2 in 2023, but everyone kind of changed the way they talk about them to say 'open-weights models' because no one releases the data they've trained their models on."

Nemotron is built differently than Llama or GPT-OSS, highlights Kari Briski. "It isn't just the models, it's the datasets, both the pre-training and post-training datasets, as well as the algorithms, software, and research. You can pick up our datasets, you can go through them, you can know what's inside of them. If you decide to pre-train your own or train your own data, you can pick and choose parts of the dataset that we've put out there – or not." There's no dearth of choice and flexibility in what Nemotron can do.

The point, for enterprises, is practical: "What that allows enterprises to do is take these open models and then blend them with their proprietary

data or their IP or their personal information that they have to train on, and so that the model can run safely on-prem and in-house – none of that data ever comes back to us," emphasizes Briski. "Having that sort of open data really allows the enterprises to have that perfect balance between proprietary data and IP, and trust – through transparency – the model enough to put those two together."

She further highlights that GenAI based agents in 2025 isn't just one model, but a system of models. "There's no one model to rule them all. On top of that, we're also saying that it takes specialized AI in different sizes of models. Nemotron comes in many sizes – Nano, Super, and Ultra – think small, medium, and large."

If 2023 was the year of tokens-as-a-flex, 2025 is fast ending on a tokens-as-a-budget note. Efficiency is the name of the game, when a hyperscaler is deploying GenAI applications at scale, where every penny and paisa counts. "The demand for tokens is going up faster than the cost per token is going down, which is why we've made these models extremely efficient," explains Briski. "The more tokens you can generate in a short time, the larger the search space for the answer. More tokens can mean more accuracy – but we also want token efficiency."

On the how, she's comfortable getting under the hood. "If I look at token efficiency – throughput over time – our previous dense models, when we did neural-architecture search and reduced size, were about 30% smaller than a typical Llama-class model at the same accuracy. We will eventually have an MoE or Mixture of Experts – that's on the roadmap – but our model is a hybrid: a Mamba state-space model as well as a transformer. That eliminates some of the compute cost even with a dense architecture," Briski explains. **■**

Read more on Digit.in:
tinyurl.com/vazdkyp9