

The “Agent” reality check

It is hard to miss the weird split-screen moment AI is having right now. On one side, the demos are dazzling and product pages are drenched in “agentic” language. On the other, companies are struggling to sell these tools at the scale the hype suggests. Once the conference theatre fades, procurement teams and business owners ask the most awkward question in tech: “So what does it do for us next quarter?”

A big part of the problem is that the hype cycle has become the product. Investors and analysts prefer clean narratives about platform shifts over modest stories regarding reduced support handle times. Buyers are not innocent either. Many leadership teams want to be seen as “AI-first” even if their data governance is stuck in 2016. The result is a market rich in promises and comparatively poor in repeatable outcomes.

That gap is glaring when the conversation moves from chat to action. Assistive tools, summarisation, and basic copilots deliver value with low risk because humans stay in charge. Agentic tools aim higher, promising to execute across systems. This is where things get messy. Hallucinations become operational hazards when an agent files the wrong ticket or messages the wrong customer. Even small error rates look unacceptable once the tool touches real workflows.

Consequently, vendors reach for guardrails. They narrow scopes, limit actions, and build elaborate monitoring. This is an admission that “autonomous” is still more of a direction than a shipped reality. When a product needs layers of constraints and constant tuning to be deployable, selling it becomes harder. This is particularly true for large enterprises that must explain decisions, audit actions, and survive regulatory microscopes.

This is where the sales story often wobbles. Security teams want to know where data goes. Legal teams want to know who is liable when the system errs. Finance teams want an ROI model more concrete than “productivity.” These demands stretch timelines and turn broad rollouts into narrow, controlled



“The problem is not that AI is a mirage. The problem is that the industry has spent so much time selling the dream that it neglected the unglamorous work of trust: reliability, governance & accountability”

Mithun Mohandas
Managing Editor



Let me know your thoughts on this column at:
@mithun_mohandas | mithun@digit.in

trials. It is not surprising that some vendors are quietly rebranding “agents” as “workflow assistance” to reduce friction.

The broader narrative about an AI bubble adds a kind of sales tax. Enterprise buyers are paid to be cautious, so organisations hedge. They sign shorter deals and delay big bets until the category feels less volatile. Ironically, the more vendors oversell inevitability, the more they encourage skepticism. Even the biggest names are feeling the high training and talent costs. At some point, the market wants businesses, not just breakthroughs.

Training, inference, and talent are expensive, and there have been repeated reports that major AI players are still spending aggressively relative to revenue. That is not a moral judgement, it is simply what happens when a technology wave is both capital-intensive and moving quickly. But it does add pressure. At some point, the market wants businesses, not just breakthroughs. When the pressure rises, optimism gets recalibrated into focus: fewer moonshot promises, more “here is the one workflow we can measurably improve”.

None of this should be read as reassurance for workers. Sales struggles today does not mean no job disruption. Capability improves in iterations. The first impact is task erosion: less writing, less research, fewer junior hours on routine analysis. Then comes role compression: teams can deliver the same output with fewer people because the “boring middle” of the work has been automated. Employers will still reshape work if the tools are “good enough” at the right slice of activity.

The problem is not that AI is a mirage. The problem is that the industry has spent so much time selling the dream that it neglected the unglamorous work of trust: reliability, governance, and accountability. The next phase will not be won by whoever shouts “agent” the loudest. It will be won by tools that are predictable, produce measurable outcomes, and fail in ways organisations can tolerate. Less theatre, more boring competence. The hype might be struggling to close deals, but the underlying progress is not taking a break. **d**