



OpenAI from 2021 to 2026

How OpenAI went from banning AI erotica to building it

—By **Vyom Ramani**

OpenAI knew exactly why it shouldn't do this. In 2021, OpenAI discovered its AI was generating violent sexual content nobody asked for — incest scenarios, unprompted erotica, material that appalled the people building it. A former employee summed up the response plainly: "We didn't want to be just an erotica company." They pulled the models, banned erotica outright, and trained that ban into ChatGPT at launch. The decision wasn't made in ignorance. It was made with complete knowledge of what the alternative looked like.

Five years later, OpenAI is trying to become precisely that. And the most damning part isn't that they're doing it. It's that they knew, better than anyone, exactly why they shouldn't.

The slow rationalisation

The ban held until the money started talking. By 2024, porn-adjacent product ideas were circulating internally. They fizzled, but their emergence signalled something had already shifted — erotica was no longer a safety problem, it was a product category being evaluated. In

August 2025, Sam Altman went on a podcast sounding genuinely conflicted, calling a sex bot a temptation resisted. In hindsight, it sounds less like a principle and more like a man talking himself into something he hadn't quite decided yet. He made the decision two months later.



The announcement itself tells you everything. Hours after OpenAI unveiled a new wellbeing advisory council which was created, in the company's own words, to "help define what healthy interactions with AI should look like for all ages," Altman posted on X that adult mode was coming in December. The council didn't know. The safety teams didn't know. A decision

committing a \$300 billion company to its most controversial product move was made on the fly and announced on social media the same day the safeguard launched.

By 2026, the reckoning was predictable. The age verification system was misclassifying minors as adults 12% of the time. Across 100 million weekly underage users, that's millions of children. The well-being council convened, unanimous and furious. One member warned OpenAI risked building a "sexy suicide coach." The launch was delayed repeatedly.

The competitive context matters too.

Grok has spent years dismantling every guardrail associated with LLMs, treating each rollback as a refinement and each boundary removal as a milestone. OpenAI is watching its market share and making calculations accordingly.

They didn't want to be just an erotica company. Then the money got loud enough. It turns out the operative word was just. tinyurl.com/81haOaz