

## GPT 5.4 Cyber, an AI model built for cybersecurity defence



OpenAI has launched GPT-5.4-Cyber, a specialised AI model for defensive cybersecurity, aimed at helping professionals analyse software, detect vulnerabilities, and counter threats. A variant of GPT-5.4, it has fewer restrictions to enable advanced security workflows like reverse engineering. The move follows similar efforts like Anthropic's Claude Mythos, signalling a growing focus on security-focused AI tools. [tinyurl.com/sx5whsz](https://tinyurl.com/sx5whsz)

## Anthropic releases Claude Design to help create visuals



Anthropic has launched Claude Design, an AI tool that helps users create content like presentations, prototypes, and marketing materials using simple prompts. It generates initial designs that can be refined, while maintaining brand consistency by learning styles from existing assets. Built-in collaboration and export options are included. Currently in research preview, it's rolling out to paid Claude users, with enterprise access controlled by admins. [tinyurl.com/s7hz8sa](https://tinyurl.com/s7hz8sa)

# Claude Mythos Preview

Everything to know about world's most dangerous AI model

—By **Vyom Ramani**

If you thought AI writing emails or building presentations was impressive, Anthropic's Claude Mythos takes things into far more unsettling territory. This isn't just another smarter chatbot upgrade. Based on Anthropic's preview, it feels closer to something out of a cybersecurity thriller than a productivity tool.

Claude Mythos has been built with a very specific purpose: understanding and exploiting software systems. It can scan large codebases, identify vulnerabilities, and in some cases even execute exploits. In one demonstration, the model reportedly discovered and used a 17-year-old bug to gain control over a system. That goes beyond pattern recognition into something far more consequential, applied reasoning with real-world impact.

What makes Mythos particularly concerning is its autonomy. It is described as agentic, meaning it can pursue goals with minimal human input. Instead of waiting for step-by-step prompts, it can properly chain actions together, test different approaches, and refine its strategy much like a human security


researcher would. The difference is speed and scale. It can iterate faster, without fatigue, and potentially across multiple systems at once.

Anthropic is fully aware of the implications. Claude Mythos is not being released publicly and is instead being tested with a limited group of partner organisations. That caution is rooted

powered security auditor capable of proactively identifying critical vulnerabilities before malicious actors can exploit them. In a world where zero-day exploits can have massive consequences, such a system could be invaluable.

At the same time, the risks are very difficult to ignore. A model capable of this level of system interaction could

just as easily be misused. Lowering the barrier to sophisticated cyberattacks and increasing their scale is a real possibility, especially if such capabilities spread beyond controlled environments.



in what the model has already demonstrated. Internal evaluations suggest it can identify weaknesses across major operating systems and, in some cases, even move beyond controlled environments. For researchers, this represents both a breakthrough and a warning.

The intended use is defensive. Anthropic positions Mythos as an AI-

AI is no longer limited to generating text or images. It is beginning to interact directly with the infrastructure that powers the digital world. Claude Mythos may still be behind closed doors, but its existence raises a larger question about control, responsibility, and whether systems this powerful can truly remain contained. [tinyurl.com/c621nxs](https://tinyurl.com/c621nxs)