

Google I/O 2026

Antigravity is putting AI on steroids these 5 ways

Vyom Ramani | editor@digit.in

Among the more surprising revelations from I/O is that Gemini CLI is being retired. When Google released Gemini CLI last year, its purpose was to make Gemini available right inside the terminal. Having garnered over 100,000 stars on GitHub, 6,000 merged pull requests, and millions of users, Google claims the tool has simply been outgrown. The issue is not about popularity but scope. Users increasingly need to employ multiple agents that can communicate with each other and divide the workload and resolve complex issues, something that requires terminal tools to operate on a common backend with everything else developers are doing.

Antigravity CLI is the new tool, and it's much more than just a rebranded version of the old one. Antigravity CLI is built in Go and is therefore more responsive and faster, offers asynchronous workflows to control multiple agents in the background, and uses the same agent harness as Antigravity 2.0, ensuring that any improvement to core agents will apply to both products at once. The migration period ends on June 18, 2026, after which Gemini CLI and Gemini Code Assist

IDE extensions will no longer process requests from Google AI Pro and Ultra customers, as well as those using the free tier.

An agent control tower

Antigravity 2.0 is the new stand-alone desktop application optimized for the agent experience. It serves as the main hub where you interact with multiple agents running parallel tasks, dynamic subagents for parallelizing tasks, scheduled tasks for background automation, and ecosystem integration from Google AI Studio, Android, and Firebase. Imagine this as not just an interface but a command center where you could run a massive code refactoring project along with some research work concurrently without worrying about locking your entire workflow.


Managed Agents

The introduction of the Managed Agents in the Gemini API means that you can create an agent that can reason, use tools, and execute code in an isolated Linux environment with one API call. Every conversation spawns an environment that can be used further in subsequent calls while retaining everything that was done previously. This is Google's

infrastructure move where it gives developers the same agent harness that it uses internally, which is co-optimized with Gemini 3.5 Flash. You can extend the agent with custom instructions and skills using markdown files as new custom agent templates are provided in the Google AI Studio Playground.

Gemini 3.5 Flash

None of this will work unless the underlying model is fast enough to handle it. Gemini 3.5 Flash beats Gemini 3.1 Pro across almost all benchmarks while being 4x faster than other frontier models, which is exactly what you want when implementing agentic workflows that involve agents waiting on each other. That's the distinction between an AI assistant and an AI workforce.



Google AI Studio

The Google AI Studio mobile application, which you can pre-register for, will enable you to capture your idea while on-the-go and create a working prototype before you even sit down at your desk. This represents a true change in the beginning of the development loop. Furthermore, agents can now make native calls to Google Workspace APIs and integrate them in their applications, and entire projects can be exported to local development in Antigravity in one click, while preserving the full project context. Moreover, there's native Android support, meaning that developers can now create top-notch Android applications using prompts alone, and publishing them straight to the Google Play Console test track from within Google AI Studio. >

