

Is AI juggernaut slowing down to make AI more safe?

Literally in a fortnight, first Anthropic then OpenAI, two of the world's leading AI companies had to postpone the launch of their latest frontier AI model – Fable 5 and Mythos 5 for Anthropic, GPT 5.6 for OpenAI. Why? The US government told them to do it.

In the case of Anthropic, they had already launched Fable 5 and Mythos 5 for paying customers of Claude all over the world, before access was revoked on June 12th. Since Anthropic is a US company, the US Department of Commerce restricted the export of Fable 5 and Mythos 5 to worldwide users. Despite Anthropic's guardrails, the US govt still deemed Mythos and Fable 5's cyber capabilities a threat in the hands of potential adversaries.

Within days of Anthropic's censure, the same fate befell OpenAI. On June 25th, the White House asked OpenAI to delay its GPT-5.6 launch in stages. As a result, the latest ChatGPT model debuted not paying customers worldwide – like every previous model launched by OpenAI – but only to a handful of government-vetted partners. OpenAI CEO Sam Altman told staff that with the US government's help they were "approving access customer by customer."

This begs the questions, is the AI juggernaut finally slowing down? And despite its own selfish interest, is the US government indirectly pushing the pause button and curbing the enthusiasm of eager technologists at leading AI companies? And to what end, for better or worse?

According to pro-safety group Public First's Brad Carson, this is not a good precedent. It's a form of gatekeeping of cutting-edge capability of frontier tech which isn't really thought through. What the US government is doing in its capacity to restrict Anthropic and OpenAI from releasing their latest AI models isn't grounded in formalised laws or regulation, the justification is too opaque and possibly illegal.



"That frontier AI models are restricted from a national security framing rather than safety governance driven by safety researchers having a voice at the table."

Jayesh Shinde
Executive Editor



Let me know your thoughts on this column at:
jayesh.shinde@digit.in | @jshinde

The US government would argue in the contrary, pointing to how it's all part of Pax Silica – a US-led initiative that came into effect from December 2025 to secure semiconductor chips and frontier AI models.

There's a philosophical argument to be made here in support of what's happening with Anthropic and OpenAI's situation. AI industry pioneers Geoffrey Hinton and Yoshua Bengio, along with 22 co-authors, have spent years trying to hit the pause button on reckless AI advancements – especially in response to worrying capabilities like what Mythos has demonstrated.

However, there's a flipside to this argument as well. Bengio warned how licensing and distribution restrictions of AI models could be detrimental. If only a handful of labs controlled the most powerful AI systems, Bengio argued, then the resulting concentration of power threatens democracy and financial markets.

What's happening now is exactly what Hinton and Bengio warned against. That frontier AI models are restricted from a national security framing rather than safety governance driven by safety researchers having a voice at the table. Even Anthropic - the company being restricted - has argued that the standard being applied by the US government is bad for the tech industry, as it "would essentially halt all new model deployments."

This has a direct impact on countries like India, and the Global South, for which access to frontier AI access has direct implications with improving productivity, efficiency, and undertaking population scale AI-enabled empowerment efforts. It's what's propelling India's renewed push for Sovereign AI compute. Where digital services and sensitive citizen data can be hosted on indigenous, locally regulated clouds rather than relying exclusively on foreign-owned servers. Hopefully soon, our AI destiny stays within our own control, and not dictated by other nations. **d**