

goal, working through multi-step projects on its own. Use it when you want to delegate a whole task — say, compile a report or plan a trip — and return to completed work rather than instructions.

<https://manus.im>

Activepieces — Open-source automation with MCP support, priced around \$5 per flow for a budget-friendly path to AI workflows. It is the cost-conscious, open alternative to the bigger platforms. Use it when you want self-hostable automation without the hefty price tags of the better-known no-code builders.

<https://www.activepieces.com>

Relay.app — A simple no-code builder for agents and workflows with genuine human-in-the-loop steps, so a person can approve before anything irreversible happens. Use it for lightweight admin automations where you want speed but still want a human checkpoint on the actions that actually matter.

<https://www.relay.app>

SHIP CODE AND APPS FASTER

Cline — A free, open-source AI agent for your terminal or IDE that you power with your own API key, getting near-premium agentic coding for just the model costs. Use it when you want a capable coding agent without a monthly subscription — you pay only for the tokens you actually burn.

<https://cline.bot>

Bolt.new — A browser-based full-stack app builder with instant

deploy, turning prompts into running apps. It is token-based at roughly \$25 a month. Great for instant prototypes — just keep an eye on token spend, which can spike during long debugging sessions when the app keeps regenerating.

<https://bolt.new>

Lovable — A natural-language full-stack builder with a Supabase backend that runs on Claude under the hood. Free tier; paid around \$25 a month. The standout choice for non-coders who want to describe a web app in plain English and watch a working, deployable version take shape.

<https://lovable.dev>

v0 (Vercel) — Turns prompts into React components using shadcn/ui, and exports standard code so you are not locked in. It is built for clean, production-ready UI. Use it to spin up polished interface components quickly, then take the code straight into your own project with minimal rework.

<https://v0.dev>

Windsurf — A budget agentic IDE offering Cursor-level features at a predictable \$15-a-month flat rate. It keeps costs simple while delivering serious AI coding help. The pick for developers who want a capable agentic editor without variable, unpredictable billing eating into the budget.

<https://windsurf.com>

Aider — An open-source terminal coding agent that is git-aware, committing changes as it works so you keep a clean history. Free, plus model costs.

Use it when you live in the command line and want an AI pair-programmer that respects version control and works the way you already do.

<https://aider.chat>

RUN AI PRIVATELY ON YOUR OWN MACHINE

LM Studio — A polished desktop GUI to browse, download and run local models, and it checks your VRAM first so you do not pick something your machine cannot handle. It is free. The friendliest way to run local LLMs visually, with no terminal commands and no data leaving your computer.

<https://lmstudio.ai>

Ollama — A command-line runtime and daemon that serves local models through an OpenAI-compatible API. Free and developer-friendly. Use it to run local AI as a service that your own scripts and tools can call, turning your machine into a private model endpoint for whatever you are building.

<https://ollama.com>

Jan — An open-source, MIT-licensed, offline ChatGPT-style assistant with no telemetry whatsoever. It is completely free. The choice for anyone who wants a familiar chat experience that is fully private and works without an internet connection — nothing is logged, nothing is sent anywhere.

<https://jan.ai>

Open WebUI — A self-hosted, ChatGPT-style web interface that sits over your local models, giving