

The great AI book-burning

The artificial intelligence industry has spent years promising to preserve, organise and democratise human knowledge. It is, therefore, grimly appropriate that its latest method of acquiring that knowledge involves slicing the spines off of rare books.

Recent reports suggest that companies supplying data to AI developers are buying vast quantities of physical books, including obscure, antique and out-of-print titles. Their bindings are removed, their spines sliced off and their loose pages fed through high-speed scanners. Once digitised, the remains may be discarded. Anthropic's own book acquisition programme, reportedly known as Project Panama, involved purchasing and destructively scanning millions of used books for training. A US court later ruled that training models on lawfully purchased books could qualify as fair use, while drawing a distinction between that and books from pirate libraries.

Legality, however, is not the same thing as wisdom. The industry is discovering that the internet is no longer the bottomless reservoir of clean human writing it once imagined. The web is increasingly filled with machine-generated articles, automated product descriptions, synthetic reviews, copied summaries and search-engine filler. Training future models indiscriminately on that material risks producing systems that learn from the output of earlier systems, gradually amplifying errors and flattening language.

Books printed before the generative AI boom offer something far more valuable. They are human-authored, carefully structured and subjected, at least in traditional publishing, to editors, proofreaders and fact-checkers. Many contain knowledge that was never properly digitised. Booksellers have reported unusual bulk demand for older and obscure titles, with data suppliers openly promoting physical books as a source of cleaner training material. In other words, AI companies have not necessarily run out of text. They are running out of text they can trust.

That makes the decision to destroy the source even harder to defend. Non-destructive



“Companies claim they are building systems to inherit humanity’s knowledge, while behaving as though the vessels carrying that knowledge are disposable”

Mithun Mohandas
Managing Editor



Let me know your thoughts on this column at:
@mithun_mohandas | mithun@digit.in

tive book scanners are not futuristic lab equipment. Libraries, archives and universities have used overhead scanners, cradle systems and specialised imaging rigs for years. They are slower, certainly, but speed is hardly a convincing excuse when the objects being processed may be irreplaceable. Not every old book is a treasure. Libraries routinely deaccession duplicate, damaged or obsolete volumes, and plenty of mass-market paperbacks can be replaced easily. But an industrial purchasing operation designed to acquire millions of books is unlikely to understand the significance of every copy passing through it.

An ordinary volume might contain an author's inscription, a translator's corrections, a previous owner's annotations or a publisher's unusual binding. An obscure regional history could be one of only a few surviving copies. Two books with identical printed text may still be distinct historical objects.

The comparison with the Library of Alexandria may sound theatrical, but the underlying anxiety is similar. A body of knowledge is being consumed in the service of power, except this time the flames have been replaced by industrial paper cutters and server racks. There is another irony. The physical book is destroyed, but its digital contents do not necessarily enter a public archive. They may instead become proprietary training data locked inside a commercial model.

AI developers could establish rarity checks, partner with libraries, fund non-destructive digitisation and release scans into accessible archives. They could also guarantee that unusual or significant copies are returned to circulation. None of these measures would prevent model development, but make it less careless. Instead, the industry wants to destroy cultural artefacts so that a chatbot might produce a better answer, or a slightly more convincing imitation of expertise.

Perhaps that is the defining contradiction of today's AI race. Companies claim they are building systems to inherit humanity's knowledge, while behaving as though the vessels carrying that knowledge are disposable. **d**