

Open-source AI is under attack

For much of the past few years, open-source AI has been a reassuring counterweight to the concentration of power in artificial intelligence. The bargain was simple: the biggest tech companies may have the money to train enormous frontier models, but researchers, start-ups and enthusiasts could still download capable open-weight alternatives, run them locally, modify them and build products without asking a giant for permission. Lately, though, there's a growing suspicion this arrangement may not last.

The latest fuel comes from an industry-wide push around AI cybersecurity. More than 100 companies, including OpenAI, Anthropic, Google, Microsoft and several cybersecurity vendors, have backed calls for stronger safeguards as AI grows more capable of discovering vulnerabilities, generating exploits and automating attacks, arguing this could sharply lower the expertise needed to attack software and infrastructure.

That is difficult to argue with. A capable model doesn't need some astonishing new zero-day vulnerability to be dangerous. Automating reconnaissance and chaining exploits at machine speed could make attacks far easier. Governments would be foolish to ignore that. But cybersecurity is also an extraordinarily convenient argument for restricting who gets access to powerful AI.

Imagine two futures. In one, powerful models stay downloadable: researchers can inspect them, companies deploy them on-premises, universities experiment without API fees, developers modify them for niche uses. In the other, models are deemed too dangerous to distribute, and access runs only through controlled APIs operated by a handful of companies able to meet strict security requirements.

Guess which future suits Big AI's business model better? Closed models hand their makers huge commercial advantages: access metered, usage priced, customers locked in, control over updates. Developers never own the underlying model and can't take it elsewhere. Open models upend that.



“We should be equally wary of a world where the organisations warning us about those risks are also the ones who profit when access to powerful AI gets centralised”

Mithun Mohandas
Managing Editor



Let me know your thoughts on this column at:
@mithun_mohandas | mithun@digit.in

Once the weights are released, somebody can optimise them, quantise them, fine-tune them, strip away unnecessary components or deploy them somewhere the original developer cannot monetise. A capable open model can become infrastructure rather than a service. That is useful for the rest of the industry and potentially inconvenient for companies spending billions building frontier models.

None of this proves AI CEOs are plotting open-source's demise in some secret meeting. The evidence is messier. Big companies still release open-weight systems. OpenAI's gpt-oss and Meta's Llama, for instance. But those releases also show how slippery “open” has become. Downloadable weights don't necessarily mean open training data, open code or unrestricted licences.

The easiest way to weaken open AI may not be to ban it outright. It could instead be to redefine what responsible AI looks like until releasing a powerful model becomes legally, financially or reputationally impractical. Governments might require tracking mechanisms, access controls or expensive evaluations before distribution. Large corporations could absorb those costs. A university laboratory or small start-up might not.

That would be unfortunate, because open models solve problems closed ones can't. Companies keep sensitive data on their own infrastructure. Researchers study model behaviour directly. Developers build specialised systems without fearing changes in their API pricing or capabilities. There are risks to releasing ever more capable models, and pretending otherwise would be naive. But we should be equally wary of a world where the organisations warning us about those risks are also the ones who profit when access to powerful AI gets centralised.

The debate should not become open versus safe. It should be about preserving both. Because if security becomes the argument that eventually puts the most capable AI behind five corporate APIs, we may discover that open-source AI was never formally killed. It was regulated into irrelevance. ■