

Bell Labs. Facebook hired him in 2013 to found FAIR, its fundamental AI research lab, giving him the budget and headcount to do at industrial scale what he'd spent thirty years doing on a shoestring. In 2018, he shared the Turing Award with Hinton and Yoshua Bengio, the field's highest honor, awarded for exactly the work everyone had told them to abandon.

The cake and the cracks

Except LeCun's next bet started forming almost immediately, and it put him at odds with the very technology his own field was about to run on. As early as 2016, he was using what became known as his cake analogy, describing intelligence as mostly self-supervised learning, with supervised learning as icing and reinforcement learning as a cherry on top, a framework meant to argue that the future of AI wasn't in labeled data or reward signals but in systems that build an internal model of the world by predicting it. It was a reasonable, even prescient point about where the field's energy should go. It also happened to describe almost nothing about how large language models work, and LeCun spent the years after ChatGPT's release saying so in increasingly blunt terms, calling autoregressive LLMs a fundamentally limited approach that would hit a ceiling no amount of scaling could fix, and arguing in public forums and on social media that a house cat has more functional understanding of the physical world than any chatbot trained purely on text.

This is the part of the arc that has no equivalent in a typical Silicon Valley story. LeCun wasn't a skeptic watching from outside. He was Meta's chief AI scientist through the exact period the company poured tens of billions of dollars into Llama, while its CEO recruited Scale AI's Alexandr Wang to run a new superintelligence lab built explicitly around scaling the

kind of models LeCun was calling a dead end. He fought about it in public, on X, with Elon Musk, with fellow researchers, in interviews where his answers were pointed enough to become their own recurring news cycle. Llama 4 landed to disappointed reviews in 2025 at almost exactly the moment Meta was restructuring its AI division around Wang's approach, and the distance between LeCun's public position and his employer's actual roadmap had become too wide to paper over as healthy internal debate.



The exit

LeCun left Meta in November 2025, after twelve years, with a LinkedIn post that credited Mark Zuckerberg for supporting FAIR and confirmed Meta would remain a partner in whatever came next, a gracious exit for someone who'd spent two years arguing against his employer's flagship strategy. What came next was AMI Labs, Advanced Machine Intelligence, founded with Alexandre LeBrun, the former head of French health tech company Nabla, and headquartered in Paris with outposts in New York, Montreal, and Singapore. The company builds world models, JEPA-style architectures that learn by predicting outcomes in an internal representation of physical reality rather than by predicting the next word in a sentence, the exact bet LeCun has been describing since that first cake slide in 2016.

In March 2026, four months after it was founded, AMI Labs closed a seed round of \$1.03 billion at a \$3.5 billion pre-money valuation, reportedly the largest seed round any European startup has ever raised. The investor list is the tell. It's not just believers, it includes NVIDIA, whose entire revenue model depends on whichever paradigm wins, alongside Samsung, Temasek, Toyota Ventures, and individual backers like Jeff Bezos and former Google CEO Eric Schmidt. That's not money betting on a research direction out of loyalty to LeCun's reputation. It's money hedging against the possibility that he's right a second time, which is a very different kind of vote of confidence, and arguably a more meaningful one.

The unfinished verdict

There's a version of this story where LeCun is simply the field's most persistent contrarian, someone whose instinct is to bet against whatever the room currently believes, and who happened to be right once. There's another version where he's the only person in this era of AI who has been saying the same thing since before it was fashionable to say it, through a boom he helped start, an employer who ignored him, and now a startup built to settle the argument with results instead of arguments. Both versions are true at the same time, and AMI Labs is the mechanism that will eventually decide which one the history books keep. The device that OpenAI is building, the models everyone else is scaling, all of it assumes LeCun's skepticism is the noise and the current LLM trajectory is the signal. LeCun has staked a Turing Award, a twelve-year career at Meta, and now over a billion dollars of other people's money on the opposite bet. He was patient enough to wait thirty years the first time. This time, he's the one holding the checkbook, which means he doesn't have to wait for anyone's permission to find out if he's right again. **d**