

The great AI price war

Chinese open models are squeezing GPT-5.6 Sol and Claude Fable 5

—By **Vyom Ramani**

The era when frontier AI research belonged almost exclusively to American laboratories is changing rapidly. OpenAI and Anthropic once set the pace and pricing standards, but July 2026 saw two Chinese labs challenge that model with systems that approach Silicon Valley's benchmark performance at dramatically lower costs, while also offering access to their weights. This is not simply China catching up. It is China changing the economics of frontier AI.

The premium tier

OpenAI's GPT-5.6 Sol and Anthropic's Claude Fable 5 sit near the top of major benchmarks. Sol has a slight edge on Artificial Analysis' Coding Agent Index, while both lead GDPval-AA v2, which evaluates real-world tasks across 44 occupations.

But these are expensive, closed models. Sol costs \$5 per million input tokens and \$30 per million output tokens, with cached inputs at \$0.50. Fable 5 operates at a similar frontier price point. Neither provides model

weights. Users pay for access, but never own the underlying model.

The challengers

Moonshot AI launched Kimi K3 on July 16, describing it as the largest open-source model released to date. It has 2.8 trillion parameters, uses a sparse mixture-of-experts architecture that

Rather than scaling, DeepSeek retained the previous architecture's 284 billion total and 13 billion active parameters, while retraining it for reasoning and agentic tasks. Terminal-Bench 2.1 performance jumped from 61.8 to 82.7, and DeepSWE improved from 7.3 to 54.4.

The cost? Just \$0.14 per million input tokens and \$0.28 per million output tokens under an MIT license.

DeepSeek's output pricing is roughly 1 percent of Sol's. That is not a minor discount. It represents an entirely different economic category.



Why this matters

These models are not merely cheaper alternatives. Kimi K3 has surpassed Claude Opus 4.8 on some benchmarks, while DeepSeek's Flash model has outperformed its larger V4-Pro version.

Kimi K3's weights are available, while DeepSeek's MIT license allows companies to download, fine-tune and deploy the model on their own infrastructure. That reduces dependence on US APIs, data-retention policies and export controls. tinyurl.com/87y2hjd

OpenAI pauses some AI training after Hugging Face incident



OpenAI has paused some AI training after identifying security risks linked to increasingly capable models, including Astra. They are strengthening security and monitoring. Astra is receiving the strictest protections after signs of "critical" cyber capabilities. OpenAI is introducing stronger isolation, tighter internet controls and multistage monitoring to detect unauthorised access. tinyurl.com/e873iu1q

Google introduces Gemini 3.7 Flash, claims it outperforms Sonnet 5



Google has launched Gemini 3.7 Flash, calling it its "most intelligent workhorse model yet" for coding and AI agents. The model improves debugging, code generation, web development and multi-step planning. It scored 43.6% on FrontierCode 1.1 and 65.3% on DeepSWE v1.1, outperforming Gemini 3.6 Flash. Google also claims it beats Claude Sonnet 5 across several benchmarks. It starts at \$0.75 per million input tokens. tinyurl.com/d8y3fer4

[SHORTS]