

# Mark Zuckerberg on AI alignment

Why no single superintelligence can be “benevolent to everyone”

—By Vyom Ramani

**M**ark Zuckerberg has released an essay where he outlines Meta’s philosophy about superintelligences, but one of the most clear-cut statements about AI alignment is hidden there. According to Zuckerberg, not even the most perfectly designed superintelligence can be benevolent towards all people. It will never happen not because of lack of technological progress, but because the concept itself doesn’t work.

Here’s why Zuckerberg says so. Most AI research institutions try to resolve the problem of alignment by designing an algorithm with such characteristics that will make it reflect the universal values. Zuckerberg completely denies this approach, claiming that “humanity is not a monoculture”. There exist multiple sets of values that contradict each other regarding issues like freedom of speech, medicine, and childcare, etc. Any AI system, regardless of its intentions, will have to take sides on the matter and, consequently, stop being benevolent for those who disagree.

Here he provides a concrete example: an AI program that would not assist a person write a letter to potential adoptive parents at a school, considering standardized testing to be ethically unacceptable. To Zuckerberg, this is not a safety victory but a case of a company imposing its values on an individual under the guise of alignment.

However, Zuckerberg’s proposed solution does not involve a smarter or better-aligned AI system. Instead, his suggestion is to create more AI systems. According to Zuckerberg, the solution to the problem of alignment is to spread superintelligence among many people instead of concentrating it in several



laboratories or government agencies. In such a scenario, when everybody had his/her own AI system aligned with his/her own goals instead of those of some company, nobody would have the power to define what was best for the society.

The idea is then taken further in what could be seen as a thought exercise: suppose there was just one super-intelligent lawyer in the world. They would win the case irrespective of whether they had a good case because of their superintelligence. Then imagine there are multiple people who have such super intelligence. Suddenly, all those inequities will disappear, and you get an equitable result. This concept is applied by Zuckerberg in cybersecurity. The world would be more vulnerable if only

some people had access to the AI that enables them to penetrate the system.

Here, Zuckerberg diverges from the gloomy rhetoric of the AI industry on how to deal with superintelligence. Whereas other industry players suggest that superintelligence needs to be controlled tightly and made accessible after testing, Zuckerberg sees restriction itself as a danger. It has always been the case throughout history that concentration of power leads to nothing good.

He’s also, quite pointedly, taking potshots at competitors. In Facebook’s new whitepaper, Zuckerberg describes how “most other labs are developing AI to serve companies, governments, or other organizations,” and how Meta’s own lab will be the lab that develops for individual humans. Alignment in Zuckerberg’s view is not about making the AI palatable. It’s about ensuring that billions of individual AIs are loyal to their respective billions of users and counting on that power dynamic to be the actual force that keeps everyone safe. How well that will work once superintelligence does come along is another issue. But it’s certainly a different approach from the one that the AI industry is currently betting on. [tinyurl.com/s13edfde](https://tinyurl.com/s13edfde)

