

Drunix debuts as NPCI's latest open source project



The National Payments Corporation of India (NPCI) has unveiled Drunix, an open source blockchain platform designed to help enterprises build and scale blockchain networks, tokenisation platforms, digital asset ecosystems and multi-organisation networks.

Built for enterprise and public infrastructure adoption, Drunix enables blockchain deployments at scale through a high-performance distributed ledger architecture. The platform is an enhanced fork of Hyperledger Fabric, the open source enterprise blockchain framework developed under the Linux Foundation.

With Drunix, NPCI aims to accelerate blockchain readiness across the ecosystem by enabling experimentation, interoperability and real-world adoption of blockchain technologies.

“Through Drunix, we aim to accelerate blockchain readiness by providing the ecosystem with foundational technologies that support experimentation, interoperability and real-world adoption. This initiative reflects our commitment to strengthening India’s digital infrastructure while contributing to the global advancement of open source innovation,” said Vishal Kanvaty, chief technology officer, NPCI.

The launch marks NPCI’s second major open source contribution to the blockchain and tokenisation ecosystem after Falcon, its platform for production-grade blockchain network management and orchestration.

Nvidia open sources DFlash for faster LLM inference

Nvidia has open sourced DFlash, a block diffusion model that speeds up autoregressive LLM inference by up to 15× while integrating with vLLM, SGLang and Hugging Face checkpoints.

Alongside the release, the team has published 20 DFlash model checkpoints on Hugging Face, deployment recipes for Blackwell and Hopper GPUs, and support for model families including Qwen, Kimi K2.6, Llama, Gemma and gpt-oss.

The release integrates natively with leading open source inference frameworks including vLLM, SGLang and the Speculators library, allowing engineers to replace

Langflow bug reveals security gaps in AI workflows

An actively exploited vulnerability in open source AI workflow platform Langflow has exposed more than 74,000 internet-facing instances, highlighting growing security risks across rapidly deployed enterprise AI infrastructure.

The flaw, tracked as CVE-2026-5027, is already being actively exploited and could allow attackers to gain full control of affected servers. The vulnerability stems from Langflow’s file upload functionality, enabling attackers to write files to arbitrary locations on a target system. Exploitation requires only a single crafted request and no credentials, as authentication is disabled by default.

The risk is amplified by the scale of exposure. Security researchers at ProCircular identified more than 74,000 Langflow instances directly accessible from the internet. With public exploit code already available, exposed deployments face an immediate threat.

The incident also highlights the dangers of insecure default configurations. Many organisations deployed Langflow without enabling authentication or implementing adequate security hardening, leaving systems vulnerable to compromise.

The issue reflects a wider trend in enterprise AI adoption. Open source AI orchestration and workflow platforms have become popular for rapid prototyping and deployment, often outside traditional IT governance processes. Security experts increasingly describe these deployments as ‘shadow AI infrastructure’ — business-critical systems operating beyond formal oversight, patch management, and access-control frameworks.