

Forty global leaders come together to launch the x402 Foundation

The operational launch of the x402 Foundation was announced on July 14, 2026, in San Francisco. The Linux Foundation acts as the parent non-profit organisation providing the vendor-neutral governance structure for the newly formed x402 Foundation, which stewards the x402 protocol, a standard designed to embed payment capabilities directly over HTTP. This allows autonomous AI agents, APIs, and applications to process financial transactions programmatically during standard web request/response cycles as they exchange data.

The x402 protocol code and framework were initially developed and contributed to the Linux Foundation by Coinbase to address machine-to-machine transaction boundaries. The standard is payment-agnostic, supporting asset types ranging from credit/debit cards to stablecoins (such as USDC) and layer-1 native cryptocurrencies (such as XRP, Solana, or Bitcoin), allowing machine-to-machine settlements at fractions of a cent.

The Foundation consists of 40 member organisations spanning traditional finance tech, cloud infrastructure, and Web3 entities. Premier members include Adyen, AWS, American Express, Circle, Cloudflare, Coinbase, Fiserv, Google, Mastercard, Monad, MoonPay, Ripple, Shopify, Solana, Stellar, Stripe and Visa, while general and associate members include networks and fintech firms such as Aleo, Fireblocks, Injective, KakaoPay, Kite AI, LayerZero Labs, NEAR, Polygon, World Liberty Financial and the Cardano Foundation.



block-parallel drafting and verification, improving GPU utilisation, reducing latency and preserving the target model's output quality.

On an eight-node NVIDIA DGX B300 system running gpt-oss-120B with TensorRT-LLM, DFlash delivered more than 15× higher throughput than autoregressive decoding at 500–600 tokens per second per user, 1.5× higher throughput than EAGLE-3, and more than 2× higher interactivity at batch size one. It also achieved up to 5.8× throughput gains in vLLM deployments and 5.1× improvements in SGLang, making it well-suited for coding, reasoning, multilingual and agentic AI workloads.

EAGLE-3 speculative decoding with a simple configuration change. Open checkpoints are also supported through TensorRT-LLM.

Unlike conventional speculative decoding, which still drafts tokens sequentially, DFlash predicts an entire block of masked future tokens in a single forward pass. The approach enables

Google open sources k8s-aibom to expose 'shadow AI' in live infrastructure

Google has open sourced k8s-aibom, a lightweight Kubernetes controller designed to detect unregistered 'shadow AI' workloads—such as models, inference servers, and agent frameworks—that developers deploy without enterprise oversight.

The tool automatically generates standard CycloneDX 1.6 machine learning bill of materials (ML-BOM) documents, creating an inventory that integrates into clusters or exports to external sinks like Google Cloud Storage. To establish an evidence trail for frameworks including the EU AI Act, NIST AI Risk Management Framework, and ISO/IEC 42001, the controller writes these ML-BOMs using a dedicated service account enforcing a DoesNotExist precondition, making records un-overwriteable once stored.

To categorise discovered assets, k8s-aibom applies a three-tier taxonomy. Declared assets are explicitly configured in workload parameters by engineers. Inferred assets represent frameworks caught dynamically via runtime signature matching. Unresolved assets denote detected AI activity where specific versions or weights cannot be verified.

Google's security engineers pushed the tool forward under the foundational thesis that 99 per cent of enterprise vulnerabilities are not in the code written in the core application, but in a very deep tree of underlying dependencies. By extending software bills of materials to machine learning weights and agentic frameworks, the tool addresses what leadership categorises as a "P0 (highest priority) challenge for the entire planet."

Operating as an unprivileged controller in its own namespace, the tool runs without the need for sidecars, privileged DaemonSets, or eBPF kernel modifications. It scans container images, environment variables, and command-line arguments to automatically identify specific infrastructure components, including runtimes (vLLM, Triton, Ollama), agent frameworks (LangChain, AutoGen, CrewAI), and vector databases (Milvus, Qdrant, pgvector).

SYSGO strengthens industrial Linux security with ELinOS 8 release

SYSGO has launched ELinOS 8, the latest generation of its industrial Linux platform engineered for secure, sovereign, and long-lifecycle mission-critical