on reputation).

When traditional software is developed using open source libraries, licence and ownership obligations are straightforward. However, AI-based systems raise a few difficult questions:

- Who owns the AI-generated output?
- Who is liable if the AI-generated output is incorrect?
- Who is accountable if the AI-generated output causes loss of reputation or economic damage?

These are questions that traditional compliance will not be able to answer.

Why licensing is only the beginning

Many people incorrectly assume that the only thing to worry about with open source artificial intelligence (AI) products is the licensing. They think if an AI tool has an open source software licence (e.g., MIT, Apache, GPL), it is worth using.

While looking for these licences is important, there are typically many other things to review before you can confidently state that your organisation will not be exposed to excessive risk by using the tool.

Many modern AI platforms have other legal or commercial obligations beyond the software licence. These include the privacy notice, acceptable use policy, terms of service, developer agreements, data handling requirements, subscription model, and so on.

A platform that seems to have low risk because of its open source licence may have many additional obligations that could significantly alter your organisation's compliance obligations. Hence, organisations need to take a much more comprehensive approach to AI governance than they did for traditional open source software.

When governance fails

The importance of governance becomes particularly clear when examining recent AI-related failures.

An AI-generated report produced for a high-profile stakeholder engagement contained significant factual inaccuracies and hallucinated information. Once distributed, the errors were quickly identified, forcing corrective action and creating both financial and reputational consequences.

Another example involved sensitive information being uploaded into a public AI system for analysis. The dataset included personal and confidential information, creating serious privacy concerns and exposing sensitive records to unnecessary risk.

Although these incidents involved different circumstances, they shared a common root cause: governance failure.

In the first case, a human review process was missing. Critical decisions relied too heavily on machine-generated output without sufficient oversight. In the second case, organisational policies governing data usage were either unclear or not properly followed.

These examples highlight an important lesson. AI failures are often not technology failures. They are policy, process, and governance failures.

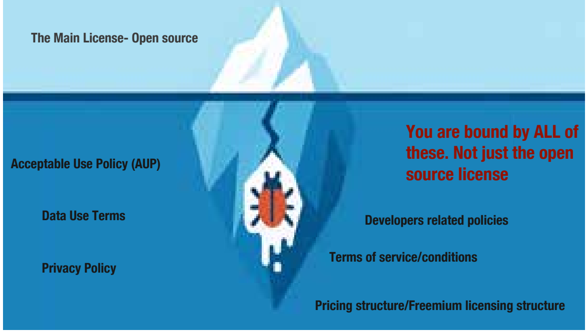

The license is just the tip of the Iceberg

Building a framework for safe innovation

Organisations need a structured framework that allows innovation while maintaining control. Across various global regulations and AI governance initiatives, four foundational pillars emerge consistently.

Governance and ethics: Every organisation using AI should establish clear governance structures. This includes accountability mechanisms, approval processes, and ethics oversight.

An AI governance board or ethics committee should evaluate high-risk use cases, define acceptable practices, and ensure that AI deployments align with organisational values and regulatory requirements.

Organisations must also understand who is using AI tools, how those tools are being deployed, and what decisions are being delegated to automated systems.

Security and risk management: The rapid growth of AI creates an equally rapid expansion of potential security risks.

Threat modelling should be incorporated into the design phase rather than be treated as an afterthought. Security reviews, red-team exercises, and continuous monitoring become essential components of responsible AI adoption.

Organisations must evaluate not only the functionality of AI systems but also their resilience against misuse, manipulation, and unintended consequences.

Data and intellectual property protection: Data remains one of the most valuable assets of any organisation. When using AI systems, companies must understand exactly what data is being shared, where that data is being processed, and how it may be used by the model provider.

The same principle applies to intellectual property. Licensing obligations must be clearly understood before AI-generated outputs are integrated into production environments.

At SAP, for example, different channels exist for experimentation, internal use, and production