


Agentic AI Needs a New Architecture



Your agents aren't broken. But the infrastructure holding them together almost certainly is.

A fraud detection agent flags a suspicious transaction and then approves it, because it has no idea that bankruptcy papers were filed for that account 40 minutes earlier. An inventory agent places a large restocking order for a warehouse that went offline an hour ago. A customer service agent routes an urgent complaint to a team that just started a scheduled break.

These are not edge cases pulled from a cautionary whitepaper. They are the kinds of failures that enterprise

architects describe, in private, when asked why their agentic AI deployments did not survive contact with production.

The agents themselves worked fine and as designed. Everything around them didn't.

According to S&P Global Market Intelligence, 42% of businesses are now scrapping most of their AI initiatives — up sharply from 17% the year before, based on a survey of more than 1,000 respondents across North America and Europe. The RAND Corporation, citing industry-wide estimates in its

2024 analysis of AI project failures, notes that more than 80% of AI projects fail — roughly twice the failure rate of comparable non-AI technology projects. In nearly every post-mortem, the model is not the primary culprit. The culprit is what surrounds it: how data flows, how systems communicate, how agents coordinate when conditions shift.

Getting this right is the difference between AI that genuinely transforms operations and AI that generates impressive demos and expensive write-offs.