

Anthropic expands Project Glasswing to 15 countries, including India



Anthropic has expanded Project Glasswing, its cybersecurity initiative developed in collaboration with open source software maintainers, to approximately 150 organisations across more than 15 countries, including India.

The expansion marks the first major rollout of the programme beyond the US and UK and grants selected organisations access to Claude Mythos Preview, Anthropic's restricted AI model for cybersecurity research, code analysis and vulnerability discovery. The company has not publicly released the model, citing concerns that malicious actors could weaponise its capabilities.

Open source organisations were among the programme's original participants, with the Linux Foundation included in the initial cohort.

"We're now expanding Project Glasswing. Following several weeks of close collaboration with our Project Glasswing partners, the security industry, open source software maintainers, and the US government, we're extending the partnership to approximately 150 new organisations," Anthropic said in a blog post.

Project Glasswing aims to identify and remediate vulnerabilities across the software ecosystems that underpin critical infrastructure. Anthropic said participating organisations have already uncovered more than 10,000 high- or critical-severity vulnerabilities using Mythos.

The company claims the model can analyse massive software codebases, uncover long-hidden security flaws and generate proof-of-concept exploit code. New participants include organisations whose services are essential to millions of people across countries such as India, Canada, Australia, France, Germany, Japan and South Korea.

Anthropic has not disclosed which Indian organisations have been granted access.

Nvidia launches open source toolkit centred on the NemoClaw framework

Nvidia has unveiled an open source Agent Toolkit designed to help developers build secure, autonomous AI workers capable of operating at enterprise scale. Announced at Nvidia GTC Taipei 2026, the toolkit combines software, open source models, open source blueprints, security frameworks, and agent orchestration tools.

At the heart of the toolkit is NemoClaw, a new open framework and blueprint for AI-agent orchestration. The framework provides templates for planning, reasoning, execution and delegation, helping developers tackle key deployment challenges such as memory management, context preservation, multi-agent collaboration and tool integration.

"NemoClaw provides enterprise software developers with the open building blocks to create more secure, long-running AI coworkers that amplify human expertise as they reshape how work gets done," said Jensen Huang, CEO of Nvidia.

The company is also addressing growing security concerns around AI agents that access sensitive files, modify code, create sub-agents and perform autonomous actions. To support secure deployment, Nvidia has introduced OpenShell Secure Runtime, a containerised environment developed with Microsoft, Canonical and Red Hat. It offers custom security controls, privacy protection, sensitive data masking, and local-only execution for critical workloads.

Nvidia has also unveiled Nemotron 3 Ultra, a 550-billion-parameter mixture-of-experts model designed for long-running autonomous agents, claiming up to 5x faster inference and 30% lower operating costs than comparable models.