Cloudflare acquires VoidZero, keeps Vite open and vendor-neutral

Cloudflare has acquired VoidZero, the open source-first company behind Vite, the world's leading JavaScript build tool, while committing to keep the Vite ecosystem open source, vendor-agnostic and community-governed.

As part of the deal, Vite creator Evan You and the core VoidZero team will join Cloudflare. The company said Vite, Vitest, Rolldown, Oxc and Vite+ will remain community-driven and available under MIT licences. Cloudflare is also committing $1 million to an independent Vite ecosystem fund that will support open source maintainers and contributors outside both Cloudflare and VoidZero, while maintaining open source steward neutrality.

The acquisition brings key open source technologies into Cloudflare, including the Vite build tool, Vitest test runner, Rolldown Rust-based bundler, Oxc toolchain, and the VoidZero engineering team.

Vite has become a foundational component of modern web development, recording more than 130 million weekly downloads. Cloudflare noted that its Vite plugin has reached 13.9 million weekly downloads, reflecting growing adoption as AI coding agents accelerate software development.

"Our mission at VoidZero has always been to eliminate the fragmentation and performance bottlenecks of the modern web stack," said Evan You, founder and CEO of VoidZero. "Joining forces allows us to keep the Vite ecosystem neutral, open, and vendor-agnostic, while giving us the resources and global infrastructure to supercharge the developer experience for millions of engineers worldwide." Cloudflare plans to integrate VoidZero's tooling into its Workers platform, enabling streamlined developer workflows, automated infrastructure provisioning and AI-native deployment experiences.

production software that developers, startups, and users can actually build with," said Paolo Ardoino, CEO of Tether.

Tether believes the technology can help move AI workloads away from centralised cloud infrastructure by enabling longer context windows and improved performance on local hardware. The company says the approach supports AI systems that process long conversations, handle large files, retain project context, support software development workflows, and work with private data locally.

Tether positions the release as a step towards more accessible, decentralised, and user-controlled AI.

Netflix engineer open sources AI cost-cutting tool

An open source tool created by Netflix senior engineer Tejas Chopra is gaining traction for tackling one of enterprise AI's fastest-growing challenges: soaring LLM usage costs.

Called Project Headroom, the software reduces AI token consumption by compressing context before it reaches a model. Since its January release, the project has reportedly saved users an estimated US$ 700,000 and recovered about 200 billion tokens for other workloads.

The need is significant. Chopra estimates that up to 90% of tokens sent to large language models can be redundant, increasing costs without improving results. By stripping unnecessary context while preserving important information, Headroom aims to lower spending and potentially improve model performance.

Headroom is fully open source and currently at version 0.22. The project has already attracted around 2,000 GitHub stars and more than 120 forks, with adoption spanning multiple Netflix teams and external projects.

"A lot of our users are people who have been really burned by token costs, more than anything else," said Chopra.

A key differentiator is its lossless, reversible compression approach. Running locally as a proxy on developer machines, the tool compresses server logs, MCP tool outputs, database outputs, file trees, documentation chunks and JSON responses. It can reportedly eliminate up to 90% of server-log data and around 70% of redundant JSON data through techniques such as CacheAligner, AST compression, JSON compression, DOM compression and Compress Cache and Retrieve (CCR).

Future plans include support for financial datasets, audio, image and video workloads, alongside a forthcoming open source project called Headlight for tracking token origins across AI workflows.

Microsoft unveils AOSP-powered platform for agent-first devices

Microsoft has unveiled Project Solara at Build 2026, positioning the Android Open Source Project (AOSP) as the foundation for a new generation of agent-first devices designed to "run AI agents instead of traditional applications."

The platform aims to replace the app-centric computing model with a task-driven approach, where AI agents interpret user intent, execute workflows across systems, and deliver results without requiring users to open individual applications.

Built as a chip-to-cloud architecture, Solara combines AOSP with Microsoft's enterprise stack, including Entra ID, Intune, Copilot-powered agents, and security and management tools. The system spans local hardware and cloud infrastructure, enabling agents to access enterprise systems, cloud data, and device inputs while coordinating tasks in the background.