

SEBI mandates ramping up cybersecurity after Anthropic AI finds 2000 vulnerabilities



Securities and Exchange Board of India (SEBI) has issued a sweeping cybersecurity directive to all regulated entities across India's securities ecosystem, after naming Anthropic's Claude Mythos Preview as a new-generation AI threat capable of discovering and exploiting software vulnerabilities at machine speed. According to the report, the AI model identified more than 2,000 previously unknown vulnerabilities within seven weeks, including flaws that had survived decades of human security review. The system also reportedly generated working exploits on the first attempt in over 83 per cent of cases.

The regulator's circular, signed on May 5 by deputy general manager Mamata Roy, applies to stock exchanges, mutual funds, depositories, credit rating agencies, portfolio managers, merchant bankers, and alternative investment funds.

SEBI warned that the interconnected nature of market infrastructure could allow a single breach to trigger widespread systemic disruption. "Due to the interconnectedness and interdependency of market participants in the Securities Market Ecosystem, a periodic coordinated approach for vulnerability management, information sharing and monitoring/assessment is required to prevent a cascading impact," the circular stated.

The directive places major emphasis on open source software supply-chain visibility. Regulated entities have been instructed to maintain updated Software Bill of Materials (SBOM) records and documentation for "all critical applications, including open source components."

SEBI also mandated immediate patching, AI-assisted vulnerability assessments, Zero Trust Network Architecture (ZTNA), API security controls, continuous monitoring, and preparations for "autonomous/agentive mitigation" against AI-driven cyber threats.

Google tests 'Remy' to compete with AI agent leader OpenClaw

Google is reportedly testing a new Gemini-powered AI agent called 'Remy', positioning the company directly against rapidly growing open source AI agent platform OpenClaw in the intensifying autonomous AI race.

According to a *Business Insider* report citing an internal document, Remy is being tested in a staff-only version of the Gemini app. The AI agent is designed to autonomously perform tasks on behalf of users rather than simply answering prompts or generating content.

The report describes Remy as a 24/7 AI assistant capable of operating across work, school, and daily life through deep integration with multiple Google services. Unlike conventional chatbots, the system focuses on proactive task automation, persistent assistance, cross-service execution, and long-duration workflows with minimal human supervision.

The project is reportedly being 'dogfooded' internally, indicating that Google employees are testing the platform before any potential public launch. No release timeline has been confirmed.

The development comes as open source AI agent platform OpenClaw gains industry attention for autonomous capabilities such as conducting research, replying to messages, and independently executing tasks.

Venture capitalist Jason Calacanis previously stated that "killing OpenClaw is the number one goal in the large language model space," while NVIDIA's Jensen Huang described the platform as "definitely the next ChatGPT."