

A key differentiator in the initiative is the integrated engineering model. Red Hat engineers will work directly within Nissan's development pipeline, allowing the automaker greater control over its software stack and faster global scaling of innovation.

Francis Chow, vice president and general manager, In-Vehicle Operating System, Red Hat, said the collaboration aims to establish "a reliable, open source foundation for the next generation of in-vehicle computing."

Google uncovers first AI-assisted zero-day exploit

Google has warned that cybercriminals and state-backed threat actors are rapidly operationalising generative AI to develop exploits, automate malware campaigns, and scale cyberattacks targeting open source infrastructure and AI ecosystems.

In a new report from *Google Threat Intelligence Group (GTIG)*, researchers disclosed the first identified real-world zero-day exploit developed with AI assistance. The exploit involved a two-factor authentication bypass targeting a popular open source web administration tool ahead of a planned mass exploitation campaign.

The operation was disrupted before deployment, following GTIG's coordinated responsible disclosure efforts with the affected vendor.

According to the report, the incident signals a major shift from experimental AI usage to operational-scale AI-driven cyberwarfare. Researchers observed sustained AI-supported vulnerability research activity linked to China- and North Korea-aligned actors, including persona-based prompting, automated exploit analysis, and agentic AI frameworks designed to scale reconnaissance and testing.

GTIG also highlighted *PROMPTSPY*, an Android backdoor capable of autonomous AI-agent behaviour through Gemini API integration. The malware can feed device interface data to Gemini, receive structured commands, autonomously click and navigate interfaces, capture biometric data, replay authentication gestures, and even prevent uninstallation attempts through invisible overlays.

Researchers additionally documented AI-assisted malware obfuscation tied to Russia-aligned operations, including dynamically generated code and AI-produced decoy logic designed to evade detection systems.

Google further warned that attackers are increasingly targeting open source AI tooling, AI integration layers, and software supply chains to facilitate credential theft, ransomware, enterprise compromise, and extortion operations. The company said it is deploying defensive AI systems, including Big Sleep, CodeMender, and expanded safeguards across Gemini services.

Kaltura open sources machine-readable AI skills

Kaltura has open sourced a suite of AI agent skills—structured, production-tested knowledge modules designed for AI coding agents such as Claude Code, OpenAI Codex, GitHub Copilot, and Cursor.

Unlike traditional documentation, these skills are machine-readable and continuously validated against production APIs. They include tested curl examples, parameter tables, error-handling logic, and security and compliance patterns, enabling AI agents to autonomously build production-grade applications.

Once deployed, the skills transform AI agents into instant platform experts, reducing development timelines from weeks to minutes and enabling end-to-end application lifecycle execution—from build and configuration to deployment and maintenance.

The capabilities span avatar-driven websites with real-time multilingual interactions, automated virtual events configured via natural language, interactive AI-powered learning environments, and fully automated content production pipelines.

Kaltura supports agent discovery and interoperability through four standards: Model Context Protocol (MCP), Agent Skills Standard, *llms.txt*, and *Context7* integration, positioning itself within an emerging AI agent infrastructure layer.

Malicious Hugging Face repository masquerading as OpenAI release hits 244K downloads

A malicious repository on Hugging Face masquerading as an OpenAI release delivered infostealer malware to Windows systems and amassed 244,000 downloads before removal, highlighting growing software supply-chain risks within open source AI ecosystems.

The fake repository, *Open-OSS/privacy-filter*, impersonated OpenAI's legitimate Privacy Filter release and copied its model card almost word-for-word to appear authentic. The repository reportedly climbed to the #1 trending position on Hugging Face within 18 hours, underscoring how public AI repositories are becoming attractive malware distribution channels.

According to AI security firm *HiddenLayer*, the repository contained a malicious *loader.py*



script that disguised itself as a legitimate AI model setup before triggering

a concealed infection chain. The malware disabled SSL verification, decoded a base64-encoded URL linked to *jsonkeeper.com*, retrieved remote payload instructions, and passed commands to PowerShell while attempting to disable Windows AMSI and ETW protections.

"Using *jsonkeeper[.]com* as the C2 channel lets the attacker rotate the payload without modifying the repository," *HiddenLayer* researchers noted.

HiddenLayer also identified six additional Hugging Face repositories using similar malicious loader logic and linked the campaign to earlier npm typosquatting attacks and fake AI packages on PyPI.