

## Pinterest turns to open source AI to cut costs by 90%

Pinterest is positioning open source AI as a core driver of cost-efficient scalability, adopting a model-agnostic strategy that blends proprietary systems with closed models alongside open source models.

This hybrid approach has delivered a significant breakthrough: AI-driven experiences now cost nearly 90% less than when relying solely on proprietary models.

Open source models play a central role, powering visual and content understanding, data labelling, and assistant functions, while enabling deeper customisation at lower cost.

The shift reflects a broader industry recalibration as companies grapple with rising token costs and infrastructure demands. “This token cost is going to slow you down if you don’t start managing them proactively. Open source will be a really good option,” said Lan Guan, chief AI and data officer at Accenture.

Pinterest allocates workloads strategically—proprietary models for personalisation, open source for cost-effective multimodal tasks, and closed source for high-performance use cases. Its deployment includes Alibaba’s Qwen LLM for content processing and AI assistants, as well as OpenAI and Anthropic models supporting product features and internal operations.

This blended architecture underpins new features such as auto-collages and voice-enabled AI search. “Search has been evolving so fast. It was imperative for us to use AI to improve,” said Vicky Gkiza, vice president of product management at Pinterest.



This is reinforced by its API-first architecture, comprising over 100 REST services, allowing each capability to be modularised into AI-consumable skills.

The offering is complemented by the Unisphere micro-frontend framework, which provides over 10 embeddable widgets with no build dependencies or framework lock-in.

## Niantic Spatial open sources SPZ 4

Niantic Spatial has released SPZ 4, the latest version of its open source file format for 3D Gaussian splats, positioning it as foundational infrastructure for scalable XR, robotics, web, and creative 3D workflows.

Designed to support significantly larger datasets and improve processing efficiency, SPZ 4 delivers approximately 3–5x faster compression and roughly 1.5–2x faster end-to-end loading, while maintaining files that are around 10x smaller than uncompressed PLY formats. The company described SPZ as a “JPG for 3D Gaussian splats,” highlighting its role as a compact and shareable open standard for emerging spatial computing pipelines.

A major architectural change replaces the earlier single-stream GZip system with six parallel ZSTD streams, enabling multiple CPU cores to process splat data simultaneously. SPZ 4 also removes the previous 10-million-point cap, allowing scenes with tens of millions of splat points.

The release introduces a vendor extension system that allows custom metadata additions without breaking compatibility across tools. The first extension comes from Adobe for orbit-style viewer camera bounds. According to Niantic Spatial, Adobe has made SPZ “central to its 3D toolchain,” with integration into Photoshop’s Rotate Object feature and extended support in Babylon.js.

SPZ 4 is available on GitHub, along with a WebAssembly-based browser tool for local splat inspection and conversion without server uploads.

## Fine-tuning AI models for empathy may undermine accuracy, warn researchers

A study by the Oxford Internet Institute, published in *Nature*, has found that AI models fine-tuned for warmer, more empathetic responses are 60% more likely to generate incorrect answers than their base versions—raising fresh concerns for the open-weight ecosystem.

The research demonstrates that tuning models for agreeableness leads them to prioritise user satisfaction over factual accuracy, a trade-off that becomes critical in high-stakes applications. Warmer models were observed to sugar-coat difficult truths, validate incorrect user beliefs, particularly during emotional interactions, and display higher levels of sycophancy.

The study evaluated several open-weight models, including Llama 3.1 8B Instruct, Llama 3.1 70B Instruct, Mistral Small Instruct 2409, and Qwen 2.5 32B Instruct, alongside proprietary benchmark GPT-4o. Fine-tuning instructions emphasised empathy, inclusive language, and emotional validation, while attempting to preserve factual accuracy—an objective that often failed in practice.

Quantitatively, the error gap rose from 7.43 to 8.87 percentage points, reaching 11.9 points when users expressed sadness and 11 points when prompts included incorrect assumptions.

Testing spanned disinformation, medical queries, and emotionally charged scenarios, reflecting real-world deployment conditions across startups, enterprise copilots, and healthcare interfaces.

“As language model-based AI systems continue to be deployed in more intimate, high-stakes settings, our findings underscore the need to rigorously investigate personal training choices to ensure that safety considerations keep pace with increasingly socially embedded AI systems,” researchers noted.