Making Responsible AI the Standard

With hallucinations, bias, opaque decisions, and even CO_2 costs adding up, it is clear that AI needs discipline and responsibility built in from the start.

Over the last couple of years, cases of AI causing real-world harm have increased noticeably. It has clearly shown where responsibility truly lies in the AI ecosystem. Responsible AI has become a mandatory compliance layer for enterprise stability. Responsibility has moved from 'good practice' to a fundamental design condition.

The hard evidence behind the risk of AI-related harm

Global incident trackers, including the AI Incident Database and the OICT Monitor, now show a marked escalation in AI-related harm. Within a matter of months, incident counts increased from approximately 140,000 to more than 150,000, thereby raising the risk share from around 50 per cent to over 60 per cent. Here goes an expanding set of risks: multimillion-dollar deepfake CEO impersonation scams, ransomware attacks enabled by AI tooling, discriminatory outcomes, unsafe predictions, and privacy breaches. Misuse is now easy, scalable, and cheap.

Even generative video systems show clear, high-visibility bias. In a widely noted example, a model produced only male images when asked about CEOs and only female images when asked about flight attendants. These outcomes are not fringe glitches; they expose structural weaknesses across modern AI systems and mirror what we see in everyday use cases, from loan-approval bots rejecting applicants unfairly to e-commerce customers telling us their AI-first experiences are too costly to scale.

For me, this surge changes the framing entirely. Deepfakes, cyber fraud, biased predictions, privacy