

### Q. How do you ensure AI safety and prevent misuse of M?

A. We use open source models but keep them fully quarantined on internet access; no memory, no data accumulation, and no data sent out. Models only process input and return output, ensuring user safety. To reduce hallucinations, we use strict domain specialisation: each query is handled by an expert model trained only on a specific Indian knowledge set. If a question is outside its domain, it won't answer. M is a mixture of focused expert models, prioritising accuracy, security, and trust unlike AGI-driven systems that aim to accumulate broad knowledge.

### Q. What is the business model behind M?

A. Our AI business is in hypergrowth and is set to surpass our 13-year-old cloud business within a year. AI is now embedded everywhere, and even riding this wave is creating strong value. Our revenue model varies by use case: in travel, we act like a travel agent and earn commissions; professional tools like coding assistants are charged hourly; while general consumer services are kept free. For legal AI, basic advice will remain free, but specialised tasks like notice or document review may be charged. On the enterprise side, corporate legal AI is already priced at ₹100 per user.

### Q. What features of M are currently available to users?

A. The mobile app offers the most complete experience today and is recommended for anyone who wants to try it. Some guardrails are intentionally in place, especially for sensitive domains like healthcare, for example, to avoid unsafe medical advice. Features will continue to evolve as safety, reliability, and user experience mature. Users can explore the mobile version now, and more capabilities will be rolled out over time.

### Q. How significant is your GPU investment?

A. We recently invested ₹36 billion in GPU infrastructure, which is expected to be fully subscribed even before deployment. GPUs are in massive demand. While use-case revenues will take time to match infrastructure revenues, we are already executing over 153 AI projects across industries, including automotive, manufacturing, and government.

### Q. How did you go about building your AI capabilities?

A. Our cloud business continues to fund the company and pays everyone's salaries, but about two years ago, we made a deliberate long-term bet on AI. We assembled a team of around 120 young engineers, invested nearly ₹500 million in GPUs, and gave them a full year to experiment internally without customer pressure. Their first work focused on applying AI to our own systems and internal enterprise chat platform, where a bot could understand and respond using three years of conversations, and a secure document repository where users could query and extract intelligence from their files. That year was about learning and capability-building, with strong support from partners such as NVIDIA, AMD, and Red Hat. Once the team was ready, we took these capabilities to customers. Over the last year alone, this has translated into more than 150 AI projects delivered.

### Q. How are you approaching sustainability and energy for data centres?

A. We are net zero even today, based on an earlier commitment, largely through renewable energy. However, this is not structurally perfect, as night-time operations still rely on grid and banked power. With 4,000 GPUs being deployed and room for

another 4,000, we will soon need a new facility. GPU density has fundamentally changed data centre design; what earlier needed 30 racks now fits into one, dramatically increasing power and cooling requirements.

For long-term sustainability, we are exploring alternative energy models, including discussions with players such as Rolls-Royce and Westinghouse on Small Modular Reactors (SMRs), with deployment targeted for around 2030. While SMRs are not yet cost-competitive, nuclear remains the only clean, consistent energy source at scale. In the near term, India's solar capacity combined with underutilised thermal power can support 9–10 GW of data centre demand; beyond 10–20GW, nuclear power becomes unavoidable.

### Q. Where is the data centre upgrade happening, and how are you addressing the talent gap for liquid cooling?

A. The GPU deployment is currently limited to a single location in Bengaluru because we are using liquid cooling, which is very different from traditional air-cooled data centres. The entire design and maintenance philosophy changes when water is taken directly to the chip. Power densities have risen sharply, and what once needed 30 racks can now fit into one making liquid cooling inevitable. India currently lacks skilled talent for this technology, as this is the first such deployment at scale.

### Q. How are you getting the talent to execute such a programme?

A. There are no academic programmes or institutes offering training, so we are building capability from scratch by learning abroad, working closely with OEMs, and training teams in-house. The first trained batch should be ready by mid-year, after which expansion to other locations will be possible. Going