


Building Production-Ready ML Systems with Open Source Tools

To prevent transition failures when moving ML models from production to deployment, development teams need to radically rethink their review workflows. They must move away from isolated code validation and adopt context-aware review systems.



Compiling a custom Linux kernel from source is a familiar rite of passage. You configure your flags, carefully select the drivers needed for your hardware architecture, run the compilation, and wait. If the compiler finishes without throwing syntax errors, you install the kernel and reboot. But if you overlooked a crucial hardware module during configuration, you are met with a catastrophic kernel panic. The code was syntactically perfect and compiled beautifully, but it lacked the necessary environmental context to run correctly on your specific machine.

Deploying a machine learning model into a production environment

often follows a strikingly similar pattern. A data scientist might spend weeks training a complex model. In their local environment—perhaps a high-end Ubuntu workstation or within a carefully curated Docker container—the model runs flawlessly. The validation metrics look fantastic, the Python scripts execute without a traceback, and the pull request is confidently opened.

From there, the continuous integration pipeline takes over. It runs standard open source linters, executes unit tests, and packages the model. Everything passes with flying colours, and the code is merged. Yet,

once that model is deployed behind a production API and exposed to the real world, unexpected issues surface. The system does not crash with a loud, easily traceable error. Instead, pipelines fail silently, API endpoints return logically flawed predictions, and the overall quality of the service degrades over time.

The root cause of these failures is rarely bad code syntax; it is a profound lack of context. To build reliable open source artificial intelligence systems, developers and DevOps engineers must look beyond standard syntax checks. We must adopt context-aware workflows to address the hidden