


Self-Hosting LLMs Using Ollama and Docker



Want to bring AI in-house but are worried about compliance issues? Ollama and Docker can help you out.

So here is something that happened to me last month in Mumbai.

I was at a client’s office, a mid-sized NBFC, and their tech head pulled me aside after the meeting. He looked tired. Said his board had been pushing him for three months to “do something with AI” but his compliance team kept blocking every single proposal. RBI rules, customer data, the usual story.

“Vishal bhai,” he said, “is there any way we can use these AI tools without my data going to America?”

Yes. There is. And that is what I want to talk about in this article.

This whole self-hosted AI thing was a niche topic until about a year ago. Now it is showing up in almost every serious

conversation I have with CTOs and CIOs. Two tools have made it possible for normal teams to actually do this without drowning in setup work — Ollama and Docker. Let me walk you through how they fit together, what works, what does not, and the things I wish someone had told me when I started.

The case for keeping AI in-house

I was sceptical about self-hosting two years ago. The hardware was painful. The models were not impressive. Setup took days. I told a few clients to just stick with the paid APIs and wait it out.

That was the right call then. It is the wrong call now.

Today’s open models are good. Really good. Llama 3.1, Mistral, Phi-3, Gemma, Qwen, DeepSeek. Some of these run on a workstation that costs less than three months of API bills