

for a busy team. A few of them even run on Mac Minis with the M-series chips. I have a colleague who runs Llama 3.2 on his laptop while travelling.

Cost is part of the story. But when I sit across the table from CTOs, here is what they tell me:

- Their data cannot leave India. Or in some cases, cannot leave a specific data centre.
- They want response times that do not depend on whether AWS has a bad day.
- They want to test ten different prompts without watching a meter tick up.
- They are tired of being held hostage by one company's pricing decisions.

These are not made-up problems. These are the things blocking AI projects from going live.

What is Ollama?

Spend a Saturday afternoon with Ollama. Seriously. It is one of those rare pieces of software where the first impression is “wait, is that really all I have to do?”

Think of Ollama as the Docker of language models. It hides the ugly parts -- the model file formats nobody understands, the graphics card driver headaches, the Python environment chaos that we have all been through.

You install it. You run Ollama and run Llama 3.1. A few minutes later you have a working AI model with an API on port 11434. That is genuinely it.

There is also this thing called a Modelfile. It looks and feels exactly like a Dockerfile. You pick a base model, add your system prompt, set things like temperature and context length, and save the whole thing as your own model. Why does this matter? Because now your prompts are config, not code. They live in Git. Your team can review them. Your auditor can see them. Your prompts are not buried inside some Python file that only one person remembers writing.

The model library covers most of what you would want. Llama, Mistral, Mixtral, Gemma 2, Phi-3, Qwen 2.5, DeepSeek. Embedding models, too, which you will need the moment you start building search.

Why Docker is part of this story

People ask me this all the time. If Ollama is already easy, why bother with Docker on top?

I had the same question. So I tried to move a setup from my laptop to a server. Then I tried to move it from one server to another. Then I tried to give the same setup to a junior developer to work on.




Figure 1: Ollama and Docker setup