

Amazon traces NPM supply chain hacks to North Korean hacker group

Amazon threat intelligence (ATI) has assessed with medium confidence that a financially motivated DPRK-linked (Democratic People's Republic of Korea) threat actor, known as Sapphire Sleet, Stardust Chollima, BlueNoroff, CageyChameleon, and Alluring Pisces, is behind a series of major node package manager (NPM) supply chain compromises.

While the March 2026 compromise of Axios (over 100 million weekly downloads) was previously linked to this group, ATI's new findings connect three earlier incidents, Typo-crypto, Debug, and Chalk, to the same operational playbook. The campaign followed a progression starting in March 2025 with Typo-crypto, where the actor placed a trojanised core.js file communicating with npmjsf.lstore. The group escalated in September 2025 by breaching Debug and Chalk, which Wiz Research found impacted roughly 1 in 10 global cloud environments within two hours, before executing the Axios breach in March 2026.

Instead of exploiting zero-day vulnerabilities, the actor obtained publisher rights by socially engineering trusted maintainers and then executed malicious payloads via automated post-install hooks. Amazon also observed the actor was using generative AI to draft and obfuscate payloads, apply execution gating to evade sandboxes, and probe AI code-auditing systems.

DeepSeek releases Harness to build coding agents around AI models

DeepSeek has launched DeepSeek Harness, an MIT-licensed open source agent framework that lets developers build coding agents around DeepSeek models. The move takes DeepSeek beyond providing foundation model access and directly into competition with Claude Code-style agent stacks.



The open source harness gives developers an alternative to proprietary agent-development frameworks, putting DeepSeek into the broader race to control the infrastructure surrounding AI agents. As agent frameworks become increasingly important to how developers build and deploy coding systems, the move could strengthen DeepSeek's position beyond model benchmarks.

Alongside, DeepSeek has upgraded its V4-Pro model to improve agent capabilities. It has also raised V4-Pro API prices, indicating that the company expects the upgraded model's stronger agent performance to justify higher model-access costs while keeping its surrounding development framework openly available.

The combination intensifies competition with Anthropic, OpenAI, and Google, particularly in AI coding and agentic systems. According to the supplied prediction-market data, Anthropic remains the frontrunner, with an 84% probability of having the best AI model by the end of September 2026.

X open sources its ranking algorithm

X is expanding its open source codebase to include the code behind its 'For You' timeline and core ranking engine. The code is available on GitHub under the Apache 2.0 licence, giving developers and researchers access to more details about how posts are ranked and displayed.

The new codebase includes the model architecture, filters, and core ranking system. It also includes the weights applied to various signals that influence which posts are

Google's HEIR will run AI models on encrypted data

Google has released HEIR, short for Homomorphic Encryption Intermediate Representation, an open source compiler and development platform for running AI models on encrypted inputs. The tool aims to make homomorphic encryption more accessible by helping developers adapt pre-trained AI models that normally process unencrypted data to work with encrypted inputs.

HEIR is part of Google's Private Computing Toolkit, which aims to reduce the expertise required to implement homomorphic encryption. The technology allows computations to be performed on encrypted data, enabling a service to process sensitive information without accessing the underlying data.

Google demonstrated HEIR through four private-inference applications compiled with the toolchain. These include a content recommendation model, a credit-card fraud detector, an intrusion detection system that analyses encrypted network packets, and a hotword detector that identifies wake words from encrypted audio recordings. The applications were evaluated using a single-threaded CPU.

Google has collaborated with universities and hardware acceleration companies in developing the project. Its research partners include Georgia Institute of Technology, Carnegie Mellon University, UC Santa Barbara, Purdue University and the University of Edinburgh. The company has also worked with acceleration partners including Belfort, Niobium, Cornami and Optalysys to research encrypted AI workloads.

The source code for the demonstration applications is available through GitHub.