

Google joins OpenROAD initiative

Google has formally joined the OpenROAD Initiative (ORI), Inc. as a principal member, which is the organisation's highest membership tier. Aaron Cunningham, a Google employee from the Hardware Toolchains Team, has been elected to the OpenROAD Initiative Governing Board to represent Google and to advise the OpenROAD Initiative (ORI) on strategic, financial and technical matters to ensure its sustainability and continued success.

Based in California, the OpenROAD Initiative Inc. is a public benefit corporation that supports an open source electronic design automation (EDA) ecosystem. ORI's sustainability is seeded and supported by the US National Science Foundation (NSF).

The OpenROAD project provides an open source, autonomous digital chip design toolchain capable of executing a complete RTL-to-GDSII flow in less than 24 hours with no human in the loop. Grounded in academic research, the project is referenced in over 500 peer-reviewed publications and has enabled thousands of students, researchers, and start-ups worldwide to design and manufacture silicon.

Nvidia's Nova driver gains Linux 7.3 support

The Nova driver, Nvidia's open source Rust-based kernel graphics driver and successor to Nouveau, is receiving further development for the upcoming Linux 7.3 merge window. The latest changes are part of ongoing work on the Linux Direct



Rendering Manager (DRM) Rust infrastructure and the Nova driver.

The DRM Rust development for Linux 7.3 involves improving the I/O infrastructure, adding a SysMem backend for shared system memory with volatile access, and introducing RegistrationData and

RegistrationGuard features to handle device resources using lifetime annotations. These changes form part of the broader effort to integrate Rust into the DRM subsystem of the Linux kernel.

In Nova, the latest changes include work to consolidate the GPU System Processor (GSP) boot process, vGPU boot process, and TLV firmware image format.

Cloudflare refreshes its community program

Cloudflare has unveiled a refreshed community program featuring two primary tracks: Ambassadors and Community Engineers.

The Ambassadors track empowers community leaders, such as campus event hosts, content creators, and peer mentors, with platform credits, marketing assets, technical resources, and online recognition. Ambassadors can serve up to two-year terms, with initial applications open through September 6, and selected members notified by October 5.



The Community Engineers track supports open source maintainers and contributors building within a company's ecosystem, including projects like

Astro, Agents SDK, EmDash, Hono, Vinext, workerd, and quiche. Unlike the Ambassadors track, Community Engineers face no maximum term limit. To support these developers, the company has committed an additional US \$1 million in grant funding over the next two years.



For more news, visit www.opensourceforu.com

Alibaba open sources Wan-Animate-2

Alibaba's Tongyi Lab (formerly Tongyi Wanxiang) released Wan-Animate-2 on August 7, 2026, an end-to-end AI character animation framework whose model weights, inference scripts, and technical paper were made publicly available under the permissive Apache 2.0 licence on GitHub and Hugging Face.

To achieve real-time streaming at 24 frames per second (at 400×720 resolution), Alibaba introduced a specialised lightweight variant called Wan-Animate-2-Lite. Unlike traditional tools that rely on intermediate skeletal pose extraction or motion encoding, Wan-Animate-2 feeds raw driving video latents directly into a redesigned Diffusion Transformer (DiT) architecture using a dual-branch system with temporal RoPE alignment. This eliminates extraction bottlenecks, preserves fine-grained dynamics like hand movements and micro-expressions, and prevents identity drift.

Beyond pose-free animation, the system features text-driven viewpoint control via a specialised Viewpoint LoRA, allowing creators to decouple the camera angle from the driving video and render identical performances from different perspectives or with custom text-prompted backgrounds. The framework natively supports multi-character driving, non-standard aspect ratios, and diverse body types without needing retraining.

To lower hardware barriers for local deployment, the release includes quantised options (INT8 and BF16), distilled speed variants, and native integration nodes for creator workflows like ComfyUI and DiffSynth-Studio.